

```

# This Python 3 environment comes with many helpful analytics
libraries installed
# It is defined by the kaggle/python Docker image:
https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
import seaborn as sns
import statsmodels.api as sm
from scipy.stats import pearsonr
from scipy import stats
from sklearn.preprocessing import MinMaxScaler

# Input data files are available in the read-only "../input/"
directory
# For example, running this (by clicking run or pressing Shift+Enter)
will list all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/)
that gets preserved as output when you create a version using "Save &
Run All"
# You can also write temporary files to /kaggle/temp/, but they won't
be saved outside of the current session

/kaggle/input/top-spotify-songs-2023/spotify-2023.csv

```

What Does It Take to Hit the Charts

Table of Contents

1. Introduction
2. Data Cleaning
3. Data Analysis & Visualization
4. How to hit the charts on Spotify

** Key findings**

- Mode, Key and BPM play a bigger role than the rest.
- "Mode: Major / Minor" plays an important role for the success of a song.

BPM matters, and not slowing down too much is the key. It has been found:

- Best range is between 110 and 90.
- Overall average is 122.
- Optimal range: 78 to 45

Introduction The dataset contains lists of the most streamed songs from the music streaming app Spotify for the year 2023. This dataset contains BPM, musical keys, mode, and other variables.

bpm: Beats per Minute, a measure of song tempo key: key of the song mode: mode of the song (Major or Minor) danceability_ %: percentage indicating how suitable the song is for dancing valence_ %: positivity of the song's musical content energy_ %: perceived energy level of the song acousticness_ %: amount of acoustic sound in the song instrumentalness_ %: amount of instrumental content in the song liveness_ %: presence of live performance elements speechiness_ %: amount of spoken words in the song

Data Cleaning Object & Scope

- Collect, clean & analyze Spotify dataset
- Identify variable of interest
- Share findings and insights

Methodology

- Collect data using "top-spotify-songs-2023"
- Wrangle data using preprocessing, cleaning, transforming and organizing data for further analysis including predictive.
- Explore data using available techniques
- Visualize data to uncover insightful discoveries

```
df = pd.read_csv('/kaggle/input/top-spotify-songs-2023/spotify-2023.csv',encoding = 'ISO8859-1')
df.tail()
```

	track_name	artist(s)_name	artist_count	\
948	My Mind & Me	Selena Gomez	1	
949	Bigger Than The Whole Sky	Taylor Swift	1	
950	A Veces (feat. Feid)	Feid, Paulo Londra	2	
951	En La De Ella	Feid, Sech, Jhayco	3	
952	Alone	Burna Boy	1	
	released_year	released_month	released_day	in_spotify_playlists
948	2022	11	3	953
949	2022	10	21	1180
950	2022	11	3	573
951	2022	10	20	1320
952	2022	11	4	782

	in_spotify_charts	streams	in_apple_playlists	...	bpm	key
mode \						
948	0	91473363	61	...	144	A
Major						
949	0	121871870	4	...	166	F#
Major						
950	0	73513683	2	...	92	C#
Major						
951	0	133895612	29	...	97	C#
Major						
952	2	96007391	27	...	90	E
Minor						

	danceability_%	valence_%	energy_%	acousticness_%
instrumentalness_% \				
948	60	24	39	57
0				
949	42	7	24	83
1				
950	80	81	67	4
0				
951	82	67	77	8
0				
952	61	32	67	15
0				

	liveness_%	speechiness_%
948	8	3
949	12	6
950	8	6
951	12	5
952	11	5

[5 rows x 24 columns]

df.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 953 entries, 0 to 952

Data columns (total 24 columns):

#	Column	Non-Null Count	Dtype
0	track_name	953 non-null	object
1	artist(s)_name	953 non-null	object
2	artist_count	953 non-null	int64
3	released_year	953 non-null	int64
4	released_month	953 non-null	int64
5	released_day	953 non-null	int64

6	in_spotify_playlists	953	non-null	int64
7	in_spotify_charts	953	non-null	int64
8	streams	953	non-null	object
9	in_apple_playlists	953	non-null	int64
10	in_apple_charts	953	non-null	int64
11	in_deezer_playlists	953	non-null	object
12	in_deezer_charts	953	non-null	int64
13	in_shazam_charts	903	non-null	object
14	bpm	953	non-null	int64
15	key	858	non-null	object
16	mode	953	non-null	object
17	danceability_%	953	non-null	int64
18	valence_%	953	non-null	int64
19	energy_%	953	non-null	int64
20	acousticness_%	953	non-null	int64
21	instrumentalness_%	953	non-null	int64
22	liveness_%	953	non-null	int64
23	speechiness_%	953	non-null	int64

dtypes: int64(17), object(7)

memory usage: 178.8+ KB

df.columns

```
Index(['track_name', 'artist(s)_name', 'artist_count',
      'released_year',
      'released_month', 'released_day', 'in_spotify_playlists',
      'in_spotify_charts', 'streams', 'in_apple_playlists',
      'in_apple_charts',
      'in_deezer_playlists', 'in_deezer_charts', 'in_shazam_charts',
      'bpm',
      'key', 'mode', 'danceability_%', 'valence_%', 'energy_%',
      'acousticness_%', 'instrumentalness_%', 'liveness_%',
      'speechiness_%'],
      dtype='object')
```

df.describe()

	artist_count	released_year	released_month	released_day	\
count	953.000000	953.000000	953.000000	953.000000	
mean	1.556139	2018.238195	6.033578	13.930745	
std	0.893044	11.116218	3.566435	9.201949	
min	1.000000	1930.000000	1.000000	1.000000	
25%	1.000000	2020.000000	3.000000	6.000000	
50%	1.000000	2022.000000	6.000000	13.000000	
75%	2.000000	2022.000000	9.000000	22.000000	
max	8.000000	2023.000000	12.000000	31.000000	

	in_spotify_playlists	in_spotify_charts	in_apple_playlists	\
count	953.000000	953.000000	953.000000	
mean	5200.124869	12.009444	67.812172	

std	7897.608990	19.575992	86.441493
min	31.000000	0.000000	0.000000
25%	875.000000	0.000000	13.000000
50%	2224.000000	3.000000	34.000000
75%	5542.000000	16.000000	88.000000
max	52898.000000	147.000000	672.000000

	in_apple_charts	in_deezer_charts	bpm	danceability_
% \				
count	953.000000	953.000000	953.000000	953.000000
mean	51.908709	2.666317	122.540399	66.96957
std	50.630241	6.035599	28.057802	14.63061
min	0.000000	0.000000	65.000000	23.00000
25%	7.000000	0.000000	100.000000	57.00000
50%	38.000000	0.000000	121.000000	69.00000
75%	87.000000	2.000000	140.000000	78.00000
max	275.000000	58.000000	206.000000	96.00000

	valence_	energy_	acousticness_	instrumentalness_
liveness_ \				
count	953.000000	953.000000	953.000000	953.000000
953.000000				
mean	51.431270	64.279119	27.057712	1.581322
18.213012				
std	23.480632	16.550526	25.996077	8.409800
13.711223				
min	4.000000	9.000000	0.000000	0.000000
3.000000				
25%	32.000000	53.000000	6.000000	0.000000
10.000000				
50%	51.000000	66.000000	18.000000	0.000000
12.000000				
75%	70.000000	77.000000	43.000000	0.000000
24.000000				
max	97.000000	97.000000	97.000000	91.000000
97.000000				

	speechiness_
count	953.000000
mean	10.131165
std	9.912888
min	2.000000

25%	4.000000
50%	6.000000
75%	11.000000
max	64.000000

Data Cleaning

#This is to check if there are any missing values in the columns?

```
nan_values = df.isna()
```

```
any_missing_values = nan_values.any().any()
```

```
any_missing_values_in_column = nan_values.any()
```

```
missing_value_count = df.isnull().sum()
```

```
for column, has_missing in any_missing_values_in_column.items():
```

```
    if has_missing:
```

```
        count = missing_value_count[column]
```

```
        print(f"-----> Column '{column}' has {count} missing values.")
```

```
print("\nMissing Values in the Entire DataFrame?")
```

```
print(any_missing_values)
```

```
print("\nMissing Values in Each Column?")
```

```
print(any_missing_values_in_column)
```

```
print("\nMissing Value Counts in Each Column:")
```

```
print(missing_value_count)
```

```
-----> Column 'in_shazam_charts' has 50 missing values.
```

```
-----> Column 'key' has 95 missing values.
```

```
Missing Values in the Entire DataFrame?
```

```
True
```

```
Missing Values in Each Column?
```

```
track_name False
```

```
artist(s)_name False
```

```
artist_count False
```

```
released_year False
```

```
released_month False
```

```
released_day False
```

```
in_spotify_playlists False
```

```
in_spotify_charts False
```

```
streams False
```

```
in_apple_playlists False
```

```
in_apple_charts False
```

```
in_deezer_playlists False
```

```

in_deezer_charts      False
in_shazam_charts      True
bpm                   False
key                   True
mode                  False
danceability_%        False
valence_%             False
energy_%              False
acousticness_%        False
instrumentalness_%    False
liveness_%            False
speechiness_%         False
dtype: bool

```

Missing Value Counts in Each Column:

```

track_name           0
artist(s)_name       0
artist_count         0
released_year        0
released_month       0
released_day         0
in_spotify_playlists 0
in_spotify_charts     0
streams              0
in_apple_playlists   0
in_apple_charts      0
in_deezer_playlists  0
in_deezer_charts     0
in_shazam_charts     50
bpm                  0
key                  95
mode                 0
danceability_%       0
valence_%            0
energy_%             0
acousticness_%       0
instrumentalness_%   0
liveness_%           0
speechiness_%        0
dtype: int64

```

based on the above results, there are many missing values in the 2 columns:

'keys' and 'in_shazams_charts'. These missing values could have happened in the data collecting process. In this case, we choose to fill all the missig values with

'Invalid'

```

df['key'] = df['key'].fillna('Invalid')
df['in_shazam_charts'] = df['in_shazam_charts'].fillna('Invalid')

```

```
# After running these, all the missing value has been replaced with "Invalid"
```

```
# Let's look at the dataset after cleaning (removing missing values)  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 953 entries, 0 to 952  
Data columns (total 24 columns):  
#   Column                                Non-Null Count  Dtype  
---  -  
0   track_name                            953 non-null    object  
1   artist(s)_name                        953 non-null    object  
2   artist_count                          953 non-null    int64  
3   released_year                        953 non-null    int64  
4   released_month                       953 non-null    int64  
5   released_day                         953 non-null    int64  
6   in_spotify_playlists                  953 non-null    int64  
7   in_spotify_charts                    953 non-null    int64  
8   streams                              953 non-null    object  
9   in_apple_playlists                   953 non-null    int64  
10  in_apple_charts                      953 non-null    int64  
11  in_deezer_playlists                  953 non-null    object  
12  in_deezer_charts                    953 non-null    int64  
13  in_shazam_charts                    953 non-null    object  
14  bpm                                  953 non-null    int64  
15  key                                  953 non-null    object  
16  mode                                953 non-null    object  
17  danceability_%                      953 non-null    int64  
18  valence_%                           953 non-null    int64  
19  energy_%                            953 non-null    int64  
20  acousticness_%                     953 non-null    int64  
21  instrumentalness_%                  953 non-null    int64  
22  liveness_%                          953 non-null    int64  
23  speechiness_%                       953 non-null    int64  
dtypes: int64(17), object(7)  
memory usage: 178.8+ KB
```

```
# here, there are 2 problems that the collected in 2 columns  
# 'streams' and 'in_deezer_playlists' turned out not to be inputted  
# with the right data type  
# They should have been inputted as a number and the datatype of  
# columns should be either  
# int64 or float64. To fix this, let's convert them into numeric data  
# types. In that process  
# replace any non-numeric data with NA  
df['streams'] = pd.to_numeric(df['streams'], errors='coerce')  
df['in_deezer_playlists'] = pd.to_numeric(df['in_deezer_playlists'],  
errors='coerce')
```



```
# Let's check the info of the dataset again
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 953 entries, 0 to 952
Data columns (total 24 columns):
#   Column                               Non-Null Count  Dtype
---  -
0   track_name                           953 non-null    object
1   artist(s)_name                       953 non-null    object
2   artist_count                         953 non-null    int64
3   released_year                       953 non-null    int64
4   released_month                      953 non-null    int64
5   released_day                        953 non-null    int64
6   in_spotify_playlists                953 non-null    int64
7   in_spotify_charts                   953 non-null    int64
8   streams                             952 non-null    float64
9   in_apple_playlists                  953 non-null    int64
10  in_apple_charts                     953 non-null    int64
11  in_deezer_playlists                  874 non-null    float64
12  in_deezer_charts                    953 non-null    int64
13  in_shazam_charts                    953 non-null    object
14  bpm                                 953 non-null    int64
15  key                                 953 non-null    object
16  mode                                953 non-null    object
17  danceability_%                      953 non-null    int64
18  valence_%                           953 non-null    int64
19  energy_%                            953 non-null    int64
20  acousticness_%                     953 non-null    int64
21  instrumentalness_%                  953 non-null    int64
22  liveness_%                          953 non-null    int64
23  speechiness_%                       953 non-null    int64
dtypes: float64(2), int64(17), object(5)
memory usage: 178.8+ KB

# The 2 columns now are back to the right data type.
# Last but not least, Let's check if there is any duplicate in the
data set
df.duplicated().sum()
# If there is no duplicate the output should be zero

0
```

Data Analysis & Visualization

```
# First, Take a look at the statistics of the data set
df.describe()
```

	artist_count	released_year	released_month	released_day	\
count	953.000000	953.000000	953.000000	953.000000	
mean	1.556139	2018.238195	6.033578	13.930745	
std	0.893044	11.116218	3.566435	9.201949	
min	1.000000	1930.000000	1.000000	1.000000	
25%	1.000000	2020.000000	3.000000	6.000000	
50%	1.000000	2022.000000	6.000000	13.000000	
75%	2.000000	2022.000000	9.000000	22.000000	
max	8.000000	2023.000000	12.000000	31.000000	
	in_spotify_playlists	in_spotify_charts	streams	\	
count	953.000000	953.000000	9.520000e+02		
mean	5200.124869	12.009444	5.141374e+08		
std	7897.608990	19.575992	5.668569e+08		
min	31.000000	0.000000	2.762000e+03		
25%	875.000000	0.000000	1.416362e+08		
50%	2224.000000	3.000000	2.905309e+08		
75%	5542.000000	16.000000	6.738690e+08		
max	52898.000000	147.000000	3.703895e+09		
	in_apple_playlists	in_apple_charts	in_deezer_playlists	\	
count	953.000000	953.000000	874.000000		
mean	67.812172	51.908709	109.740275		
std	86.441493	50.630241	178.811406		
min	0.000000	0.000000	0.000000		
25%	13.000000	7.000000	12.000000		
50%	34.000000	38.000000	36.500000		
75%	88.000000	87.000000	110.000000		
max	672.000000	275.000000	974.000000		
	in_deezer_charts	bpm	danceability_%	valence_%	
energy_%	\				
count	953.000000	953.000000	953.000000	953.000000	
953.000000					
mean	2.666317	122.540399	66.96957	51.431270	
64.279119					
std	6.035599	28.057802	14.63061	23.480632	
16.550526					
min	0.000000	65.000000	23.000000	4.000000	
9.000000					
25%	0.000000	100.000000	57.000000	32.000000	
53.000000					
50%	0.000000	121.000000	69.000000	51.000000	
66.000000					
75%	2.000000	140.000000	78.000000	70.000000	
77.000000					
max	58.000000	206.000000	96.000000	97.000000	
97.000000					
	acousticness_%	instrumentalness_%	liveness_%	speechiness_%	

count	953.000000	953.000000	953.000000	953.000000
mean	27.057712	1.581322	18.213012	10.131165
std	25.996077	8.409800	13.711223	9.912888
min	0.000000	0.000000	3.000000	2.000000
25%	6.000000	0.000000	10.000000	4.000000
50%	18.000000	0.000000	12.000000	6.000000
75%	43.000000	0.000000	24.000000	11.000000
max	97.000000	91.000000	97.000000	64.000000

Release Date

- The year of release of the latest song in the dataset is (of course) 2023, with the oldest dating back to 1930.

Popularity

- Some of the most streamed songs never landed in the charts.
- The most number of streams is 3,703,895,000, while the least number of streams is 2,762.
- The mean number of streams is 514,137,400 with a standard deviation of 566,856,900, which indicates that the number of streams varies by a huge value.

Song Properties / Characteristics

- The highest and lowest BPM is 206 and 65, respectively, with a mean value of 122.55.
- The highest and lowest danceability is 96% and 23%, respectively, with a mean value of 66.98%.
- The highest and lowest valence is 97% and 4%, respectively, with a mean value of 51.41%.
- The highest and lowest energy is 97% and 9%, respectively, with a mean value of 64.27%.
- The highest and lowest acousticness is 97% and 0%, respectively, with a mean value of 27.08%.
- The highest and lowest instrumentalness is 91% and 0%, respectively, with a mean value of 1.58%.
- The highest and lowest liveness is 97% and 3%, respectively, with a mean value of 18.21%.
- The highest and lowest speechiness is 64% and 2%, respectively, with a mean value of 10.14%.

How to hit the charts on Spotify

```
## There are multiple relatable factors on the table that can be used
to analyze and find out what help a song hit the charts.
# We need to set the targetted dataframe apart from the original
database.
# Possible factors can be list as: 'bpm','key','mode', 'danceability_
```

```

%', 'valence_%', 'energy_%', 'acousticness_%', 'instrumentalness_%',
'liveness_%', 'speechiness_%'
# So we need to create a sub dataframe with just these columns to
examine the numbers and data
# First, we may want to change the unnecessarily long track names
into only tracks ID
track_name_to_id = {name: idx for idx, name in
enumerate(df['track_name'].unique())}
df['track_id'] = df['track_name'].map(track_name_to_id) # map the
track names to the track IDs
df.drop(columns=['track_name'], inplace=True) # replace the track
names with the track IDs

# create a list storing all the names of new columns in the sub
dataframe
voi_col = ['track_id', 'streams', 'bpm', 'key', 'mode', 'danceability_
%', 'valence_%', 'energy_%', 'acousticness_%', 'instrumentalness_%',
'liveness_%', 'speechiness_%']
# create a new sub dataframe with only columns matched with above
lists
df_voi = df[voi_col]
df_voi

```

	track_id	streams	bpm	key	mode	danceability_%	valence_%
\							
0	0	141381703.0	125	B	Major	80	89
1	1	133716286.0	92	C#	Major	71	61
2	2	140003974.0	138	F	Major	51	32
3	3	800840817.0	170	A	Major	55	58
4	4	303236322.0	144	A	Minor	65	23
...
948	938	91473363.0	144	A	Major	60	24
949	939	121871870.0	166	F#	Major	42	7
950	940	73513683.0	92	C#	Major	80	81
951	941	133895612.0	97	C#	Major	82	67
952	942	96007391.0	90	E	Minor	61	32
	energy_%	acousticness_%	instrumentalness_%	liveness_%			
	speechiness_%						
0	83	31	0	8			

4				
1	74	7	0	10
4				
2	53	17	0	31
6				
3	72	11	0	11
15				
4	80	14	63	11
6				
..
...				
948	39	57	0	8
3				
949	24	83	1	12
6				
950	67	4	0	8
6				
951	77	8	0	12
5				
952	67	15	0	11
5				

[953 rows x 12 columns]

```
# Since key and mode contain non-numeric data so we need to modify
these columns
# for analyzing. One of the ways to solve this is to change these into
binary numbers
# by utilizing the get_dummies function.
```

```
#label encoding or one hot encoding?
```

```
df_features = df_voi
df_features = pd.get_dummies(df_voi, columns=['key', 'mode'],
prefix=['key', 'mode'])
```

```
#rename columns of interests
```

```
re_col = {
    'danceability_%': 'danceability',
    'valence_%': 'valence',
    'energy_%': 'energy',
    'acousticness_%': 'acousticness',
    'instrumentalness_%': 'instrumentalness',
    'liveness_%': 'liveness',
    'speechiness_%': 'speechiness'
}
```

```
df_features.rename(columns=re_col, inplace=True)
df_features.head()
```

	track_id	streams	bpm	danceability	valence	energy	acousticness	\
0	0	141381703.0	125	80	89	83	31	
1	1	133716286.0	92	71	61	74	7	
2	2	140003974.0	138	51	32	53	17	
3	3	800840817.0	170	55	58	72	11	
4	4	303236322.0	144	65	23	80	14	

	instrumentalness	liveness	speechiness	...	key_D	key_D#	key_E	key_F	\
0	0	8	4	...	False	False	False	False	
1	0	10	4	...	False	False	False	False	
2	0	31	6	...	False	False	False	True	
3	0	11	15	...	False	False	False	False	
4	63	11	6	...	False	False	False	False	

	key_F#	key_G	key_G#	key_Invalid	mode_Major	mode_Minor
0	False	False	False	False	True	False
1	False	False	False	False	True	False
2	False	False	False	False	True	False
3	False	False	False	False	True	False
4	False	False	False	False	False	True

[5 rows x 24 columns]

Here since the track_id column is not very useful, so we decided to drop it.

```
df_da = df_features.drop(columns=['track_id'])
df_da.head()
```

	streams	bpm	danceability	valence	energy	acousticness	\
0	141381703.0	125	80	89	83	31	
1	133716286.0	92	71	61	74	7	
2	140003974.0	138	51	32	53	17	
3	800840817.0	170	55	58	72	11	
4	303236322.0	144	65	23	80	14	

	instrumentalness	liveness	speechiness	key_A	...	key_D	key_D#	key_E	\
0	0	8	4	False	...	False	False	False	

False							
1	0	10	4	False	...	False	False
False							
2	0	31	6	False	...	False	False
False							
3	0	11	15	True	...	False	False
False							
4	63	11	6	True	...	False	False
False							

	key_F	key_F#	key_G	key_G#	key_Invalid	mode_Major	mode_Minor
0	False	False	False	False	False	True	False
1	False	False	False	False	False	True	False
2	True	False	False	False	False	True	False
3	False	False	False	False	False	True	False
4	False	False	False	False	False	False	True

[5 rows x 23 columns]

```
# Binary columns contain all the keys and one of the 2 mode (either
major or minor)
# used in composing those songs on the charts. Since, we aim to
discover what to get a hit on top chart
# Figuring out facts about keys and modes may help. Now, with the
get_dummies function,
# all the data in these columns is either True or False. Let's convert
them into binary data,
# which contain either 1 or 0 value.
```

```
binary_columns = ['key_A', 'key_A#', 'key_B', 'key_C#', 'key_D',
'key_D#', 'key_E', 'key_F', 'key_F#', 'key_G', 'key_G#', 'mode_Major',
'mode_Minor']
```

```
for column in binary_columns:
    df_da[column] = df_da[column].astype(int)
print(df_da.head())
```

	streams	bpm	danceability	valence	energy	acousticness	\
0	141381703.0	125	80	89	83	31	
1	133716286.0	92	71	61	74	7	
2	140003974.0	138	51	32	53	17	
3	800840817.0	170	55	58	72	11	
4	303236322.0	144	65	23	80	14	

	instrumentalness	liveness	speechiness	key_A	...	key_D	key_D#
key_E \							
0	0	8	4	0	...	0	0
0							
1	0	10	4	0	...	0	0
0							


```
4    0.0    0.0    0.0    0.0    0.0    0.0    1.0
```

```
[5 rows x 23 columns]
```

By using the MinMaxScaler in SK learning library, we can normalize the input features

```
scaler = MinMaxScaler()
```

```
df_da['streams'] = scaler.fit_transform(df_da[['streams']])
```

```
print(df_da.head())
```

	streams	bpm	danceability	valence	energy	acousticness	\
0	0.038170	125.0	80.0	89.0	83.0	31.0	
1	0.036101	92.0	71.0	61.0	74.0	7.0	
2	0.037798	138.0	51.0	32.0	53.0	17.0	
3	0.216215	170.0	55.0	58.0	72.0	11.0	
4	0.081869	144.0	65.0	23.0	80.0	14.0	

	instrumentalness	liveness	speechiness	key_A	...	key_D	key_D#
key_E \							
0	0.0	8.0	4.0	0.0	...	0.0	0.0
0.0							
1	0.0	10.0	4.0	0.0	...	0.0	0.0
0.0							
2	0.0	31.0	6.0	0.0	...	0.0	0.0
0.0							
3	0.0	11.0	15.0	1.0	...	0.0	0.0
0.0							
4	63.0	11.0	6.0	1.0	...	0.0	0.0
0.0							

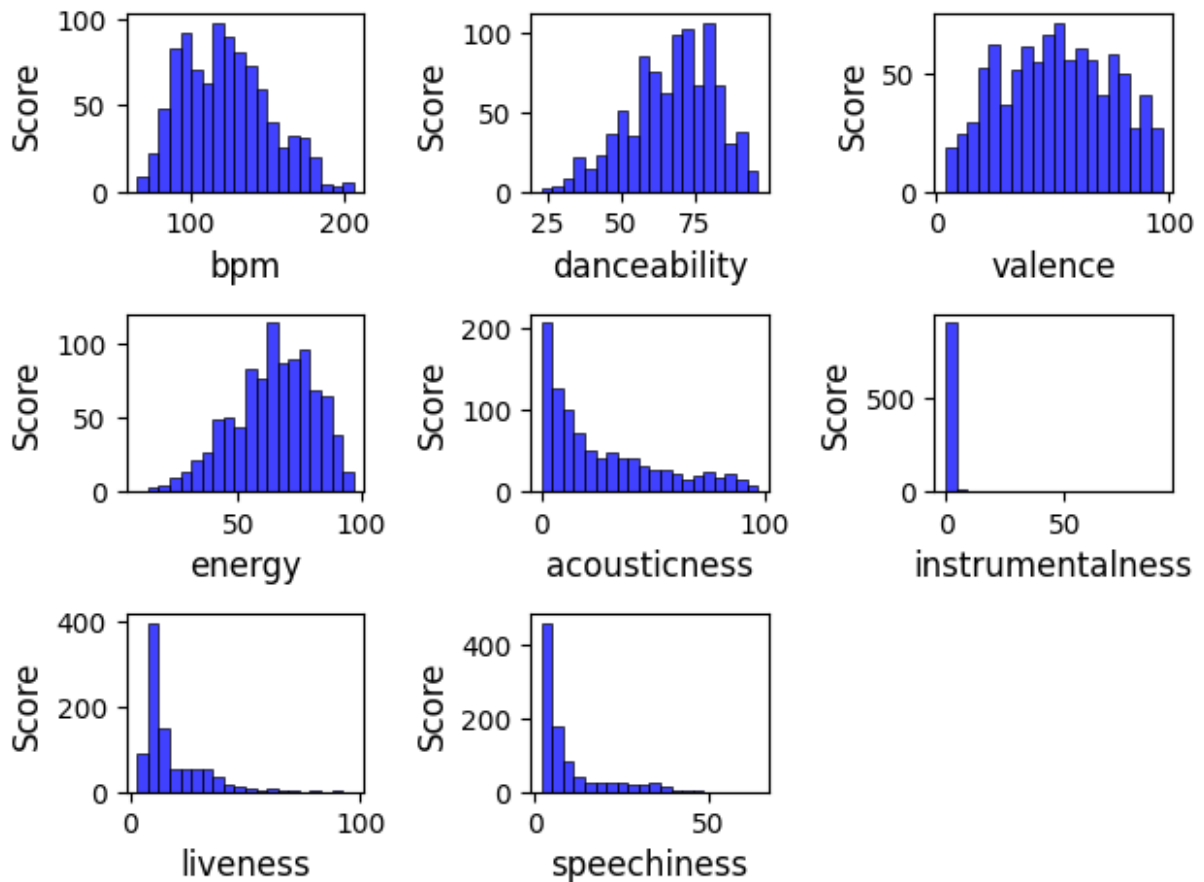
	key_F	key_F#	key_G	key_G#	key_Invalid	mode_Major	mode_Minor
0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
1	0.0	0.0	0.0	0.0	0.0	1.0	0.0
2	1.0	0.0	0.0	0.0	0.0	1.0	0.0
3	0.0	0.0	0.0	0.0	0.0	1.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0	1.0

```
[5 rows x 23 columns]
```

```
columns_to_plot = ['bpm', 'danceability', 'valence', 'energy',  
'acousticness', 'instrumentalness', 'liveness', 'speechiness']
```

```
for i, column in enumerate(columns_to_plot, 1):  
    plt.subplot(3, 3, i)  
    sns.histplot(data=df_da, x=column, bins=20, color='blue')  
    plt.xlabel(column, fontsize=12)  
    plt.ylabel("Score", fontsize=12)
```

```
plt.tight_layout()
plt.show()
```



The chart illustrates the distribution of various music variables among the top 1000 most popular songs on Spotify in 2023. It comprises a set of frequency charts, with the x-axis representing music variables and the y-axis representing the frequency of songs with those values.

The most popular songs tend to have high danceability. This might be due to their upbeat and enjoyable nature, making them more appealing to listeners.

Additionally, the most popular songs often exhibit moderate energy levels, which could contribute to their listenability and appeal.

Danceability

The distribution of danceability resembles a bell curve, indicating that most songs have a moderate level of danceability around 70%. Some songs, however, have very high danceability exceeding 90%, typically falling within the dance or pop genres.

Energy

Similar to danceability, the distribution of energy also forms a bell curve, signifying that the majority of songs exhibit a moderate level of energy, approximately 60%. Some songs have higher energy levels surpassing 80%, often belonging to the rock or pop genres.

Valence

The valence distribution also resembles a bell curve, suggesting that most songs have a moderate valence around 60%. There are songs with higher valence exceeding 80%, commonly found in pop or hip-hop genres.

Acousticness

The distribution of acousticness skews towards the right, indicating a higher presence of songs with higher acousticness compared to lower acousticness. On average, songs have an acousticness level of around 30%. Some songs surpass 60% acousticness, typically belonging to folk or indie genres.

Instrumentalness

The distribution of instrumentalness forms a bell curve, showcasing that most songs have a moderate instrumentalness level around 50%. There are songs with higher instrumentalness surpassing 70%, often comprising instrumental or classical music.

Liveness

The liveness distribution also resembles a bell curve, suggesting that most songs have a moderate liveness level around 50%. Some songs exhibit higher liveness exceeding 70%, usually associated with rock or pop genres.

Speechiness

The speechiness distribution skews towards the left, indicating a higher presence of songs with lower speechiness compared to higher speechiness. On average, songs have a speechiness level of about 20%. Some songs exceed 40% speechiness, commonly found in rap or hip-hop genres.

Overall, the chart indicates that the most popular songs on Spotify tend to have high danceability and moderate energy levels. These songs also typically exhibit moderate valence, higher acousticness, moderate instrumentalness, moderate liveness, and lower speechiness.

```
# One of the best way to figure out the correlation and relationship
# in the dataframe. We use corr() function to create the correlation
# matrix between all the columns
columns_to_correlate = ['bpm', 'danceability', 'valence', 'energy',
                        'acousticness', 'instrumentalness', 'liveness', 'speechiness',
                        'streams']
```

```
correlation_matrix = df_da[columns_to_correlate].corr()
print(correlation_matrix)
```

```
          bpm  danceability  valence  energy
acousticness \
```


bpm	1.000000	-0.147095	0.041195	0.025794	-
0.017694					
danceability	-0.147095	1.000000	0.408451	0.198095	-
0.236165					
valence	0.041195	0.408451	1.000000	0.357612	-
0.081907					
energy	0.025794	0.198095	0.357612	1.000000	-
0.577344					
acousticness	-0.017694	-0.236165	-0.081907	-0.577344	
1.000000					
instrumentalness	-0.001195	-0.089138	-0.132890	-0.038547	
0.042796					
liveness	-0.000761	-0.077538	0.021278	0.117302	-
0.050142					
speechiness	0.039260	0.184977	0.041081	-0.004846	-
0.022501					
streams	-0.002438	-0.105457	-0.040831	-0.026051	-
0.004485					

	instrumentalness	liveness	speechiness	streams
bpm	-0.001195	-0.000761	0.039260	-0.002438
danceability	-0.089138	-0.077538	0.184977	-0.105457
valence	-0.132890	0.021278	0.041081	-0.040831
energy	-0.038547	0.117302	-0.004846	-0.026051
acousticness	0.042796	-0.050142	-0.022501	-0.004485
instrumentalness	1.000000	-0.045967	-0.083396	-0.044902
liveness	-0.045967	1.000000	-0.022525	-0.048337
speechiness	-0.083396	-0.022525	1.000000	-0.112333
streams	-0.044902	-0.048337	-0.112333	1.000000

Next, take the correlation matrix between all the feature to draw a heat map

```
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
fmt=".2f", square=True)
```

```
plt.title("Correlation Heatmap")
plt.show()
```



The correlation heatmap illustrates the relationship between various music variables among the top 1000 most popular songs on Spotify in 2023. The heatmap employs a color scale to represent the strength of correlation, where shades of blue indicate a positive correlation and shades of red denote a negative correlation.

The most strongly correlated music variables are danceability and valence. This might be due to songs with high danceability often possessing high valence as well. For instance, dance songs commonly exhibit both of these characteristics.

Conversely, the weakest correlation exists between Acousticness and Energy. This could be because high Acousticness doesn't necessarily imply high energy. For example, acoustic songs might have low energy.

Danceability and Valence

The correlation between danceability and valence is very strong, with a correlation coefficient of 0.4. This suggests that songs with high danceability tend to have higher valence compared to

songs with low danceability. This correlation might be due to the lively and enjoyable melodies often associated with songs featuring high danceability, resulting in higher valence.

Energy and Valence

The correlation between energy and valence is moderate, with a correlation coefficient of 0.36. This indicates that songs with high energy necessarily have high valence. This weak correlation could be attributed to the fact that songs with high energy might have strong and lively beats necessarily leading to high valence.

Acousticness and Energy

The correlation between danceability and energy is weak, with a correlation coefficient of -0.58. This implies that songs with high Acousticness are not likely to have energy. This correlation could be explained by the fact that songs with Acousticness often feature slow and sad beats, leading to energy levels low.

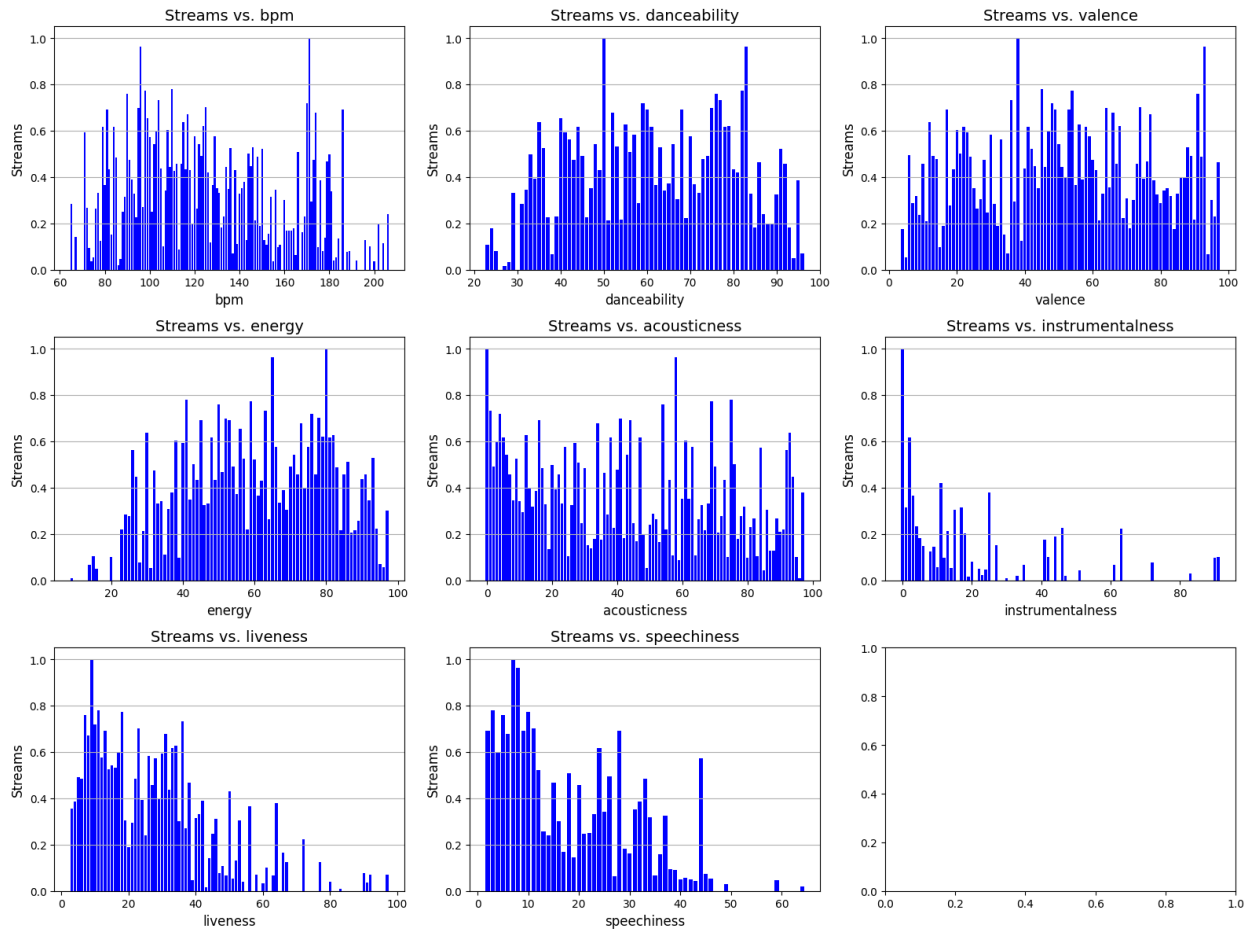
In summary, the correlation heatmap demonstrates that the strongest correlation exists between danceability and energy, while the weakest correlation is between acousticness and Energy.

```
columns_to_plot = ['bpm', 'danceability', 'valence', 'energy',
                   'acousticness', 'instrumentalness', 'liveness', 'speechiness']

fig, axes = plt.subplots(3, 3, figsize=(16, 12))

axes = axes.flatten()

for i, column in enumerate(columns_to_plot):
    plt.sca(axes[i])
    plt.bar(df_da[column], df_da['streams'], color='blue')
    plt.xlabel(column, fontsize=12)
    plt.ylabel('Streams', fontsize=12)
    plt.title(f'Streams vs. {column}', fontsize=14)
    plt.grid(axis='y')
plt.tight_layout()
plt.show()
```



The chart above shows six scatter plots that reveal various relationships between different audio features of music and the number of times those songs were streamed on Spotify.

The x-axis on each plot represents the value of a specific audio feature, while the y-axis shows the average number of streams a song gets per unit of that feature. Each data point on the plot represents a single song.

Streams vs. bpm There appears to be a weak positive correlation between the tempo (beats per minute) of a song and the number of times it's streamed. Songs with higher tempos tend to be streamed slightly more often than songs with lower tempos.

Streams vs. danceability There's a moderate positive correlation between a song's danceability and its number of streams. Danceable songs, as determined by Spotify's algorithm, tend to be streamed more often than less danceable songs.

Streams vs. valence There's a weak positive correlation between a song's valence (positiveness) and its number of streams. Songs with higher valence scores tend to be streamed slightly more often than songs with lower valence scores.

Streams vs. energy There's a moderate positive correlation between a song's energy level and its number of streams. More energetic songs tend to be streamed more often than less energetic songs.

Streams vs. acousticness There's a weak negative correlation between a song's acousticness and its number of streams. Less acoustic songs (more electronic or produced) tend to be streamed slightly more often than more acoustic songs.

Streams vs. instrumentalness There's a weak negative correlation between a song's instrumentalness and its number of streams. Songs with vocals tend to be streamed slightly more often than songs without vocals.

It's important to note that these are just correlations, and they don't necessarily mean that one feature causes another. There could be other factors that influence how often a song is streamed, such as the artist, genre, or release date.

```
columns_to_plot = ['danceability', 'valence', 'energy',
                   'acousticness', 'instrumentalness', 'liveness', 'speechiness']

fig, axes = plt.subplots(3, 3, figsize=(16, 12))
axes = axes.flatten()
colormap = plt.cm.get_cmap('viridis')

for i, column in enumerate(columns_to_plot):
    plt.sca(axes[i])

    scatter = plt.scatter(df_da[column], df_da['streams'],
                          c=df_da[column], cmap=colormap, alpha=0.5)

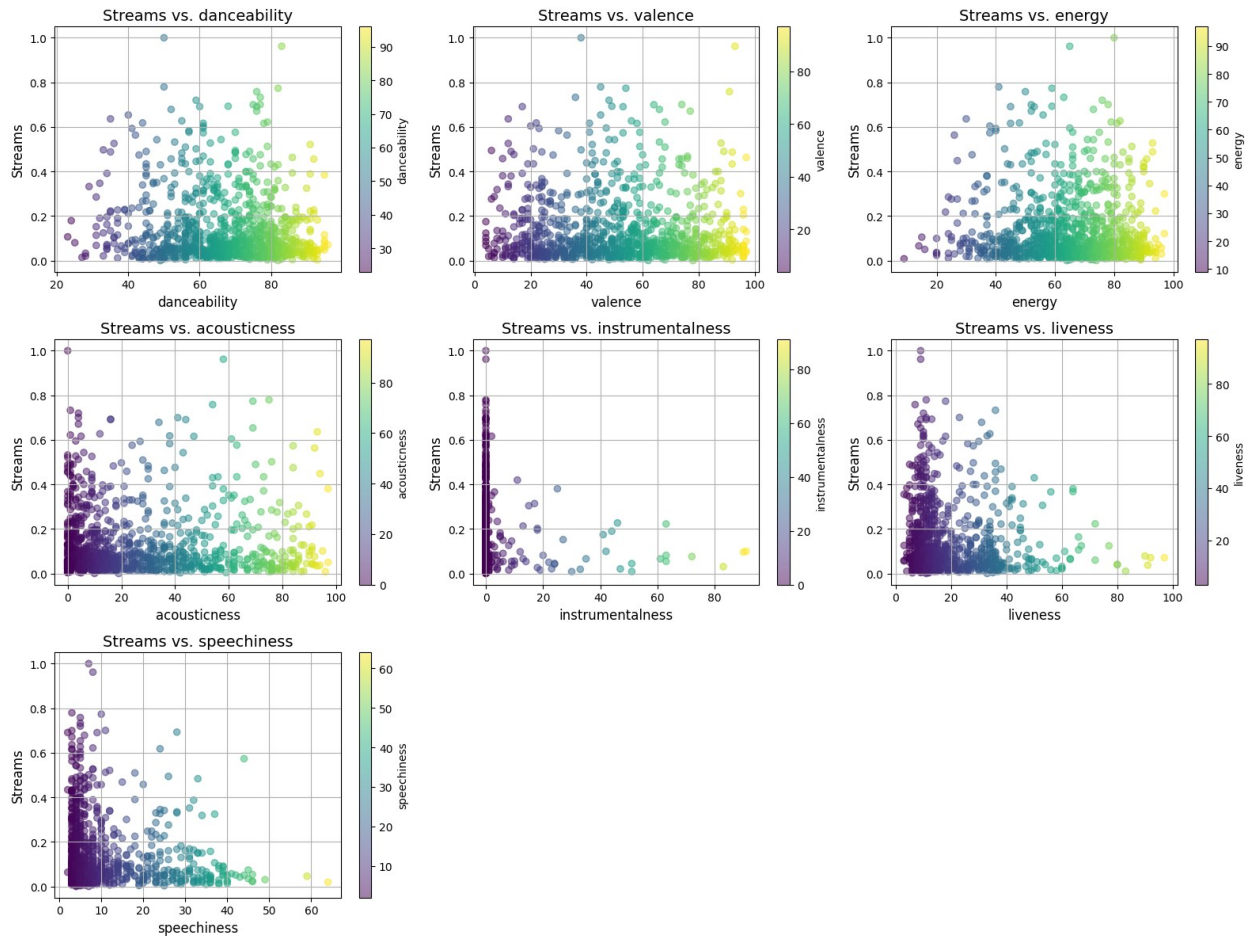
    plt.xlabel(column, fontsize=12)
    plt.ylabel('Streams', fontsize=12)
    plt.title(f'Streams vs. {column}', fontsize=14)
    plt.grid(True)
    plt.colorbar(scatter, label=column)

for j in range(len(columns_to_plot), len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()
```

/tmp/ipykernel_19/646345061.py:5: MatplotlibDeprecationWarning: The get_cmap function was deprecated in Matplotlib 3.7 and will be removed two minor releases later. Use ``matplotlib.colormaps[name]`` or ``matplotlib.colormaps.get_cmap(obj)`` instead.

```
colormap = plt.cm.get_cmap('viridis')
```

Streams vs. Danceability There appears to be a moderate positive correlation between a song's danceability and the number of times it's streamed. This means that songs that Spotify's algorithm identifies as more danceable tend to be streamed more often than songs that are less danceable. The data points are somewhat spread out, so there are definitely exceptions to this trend, but the overall pattern is clear.

Streams vs. Valence There's a weak positive correlation between a song's valence (positiveness) and its number of streams. This means that songs with higher valence scores tend to be streamed slightly more often than songs with lower valence scores. However, the spread of the data points is quite large, so there are many songs that don't follow this pattern.

Streams vs. Energy There's a moderate positive correlation between a song's energy level and its number of streams. This means that more energetic songs tend to be streamed more often than less energetic songs. Similar to the danceability plot, the data points are somewhat spread out, but there's a clear overall trend.

Streams vs. Acousticness There's a weak negative correlation between a song's acousticness and its number of streams. This means that less acoustic songs (more electronic or produced) tend to be streamed slightly more often than more acoustic songs. Again, the spread of the data points is quite large, so there are many exceptions to this trend.

Streams vs. Instrumentalness There's a weak negative correlation between a song's instrumentalness and its number of streams. This means that songs with vocals tend to be streamed slightly more often than songs without vocals. However, the spread of the data points is quite large, so there are many songs that don't follow this pattern.

Streams vs. Speechiness There is no clear correlation between a song's speechiness and the number of times it's streamed. The data points are spread evenly across the chart, with no discernible pattern.

It's important to note that these are just correlations, and they don't necessarily mean that one feature causes another. There could be other factors that influence how often a song is streamed, such as the artist, genre, or release date.

```
df_da = df_da.merge(df[['key']], left_index=True, right_index=True,
how='inner')
df_da.head()
```

	streams	bpm	danceability	valence	energy	acousticness	\
0	0.038170	125.0	80.0	89.0	83.0	31.0	
1	0.036101	92.0	71.0	61.0	74.0	7.0	
2	0.037798	138.0	51.0	32.0	53.0	17.0	
3	0.216215	170.0	55.0	58.0	72.0	11.0	
4	0.081869	144.0	65.0	23.0	80.0	14.0	

	instrumentalness	liveness	speechiness	key_A	...	key_D#	key_E
key_F \							
0	0.0	8.0	4.0	0.0	...	0.0	0.0
0.0							
1	0.0	10.0	4.0	0.0	...	0.0	0.0
0.0							
2	0.0	31.0	6.0	0.0	...	0.0	0.0
1.0							
3	0.0	11.0	15.0	1.0	...	0.0	0.0
0.0							
4	63.0	11.0	6.0	1.0	...	0.0	0.0
0.0							

	key_F#	key_G	key_G#	key_Invalid	mode_Major	mode_Minor	key
0	0.0	0.0	0.0	0.0	1.0	0.0	B
1	0.0	0.0	0.0	0.0	1.0	0.0	C#
2	0.0	0.0	0.0	0.0	1.0	0.0	F
3	0.0	0.0	0.0	0.0	1.0	0.0	A
4	0.0	0.0	0.0	0.0	0.0	1.0	A

[5 rows x 24 columns]


```
palette = sns.color_palette("Set2", len(df_da['key'].unique()))
```

```
sns.catplot(y="key", x="streams", data=df_da, aspect=2,
palette=palette)
plt.xlabel("Streams", fontsize=20)
```

```
plt.ylabel("Keys", fontsize=20)
plt.show()
```

/tmp/ipykernel_19/1760608335.py:3: FutureWarning: Passing `palette` without assigning `hue` is deprecated.

```
sns.catplot(y="key", x="streams", data=df_da, aspect=2,
palette=palette)
```



```
palette = sns.color_palette("Set2", len(df_da['key'].unique()))
```

```
sorted_keys = sorted(df_da['key'].unique())
```

```
plt.figure(figsize=(12, 6))
```

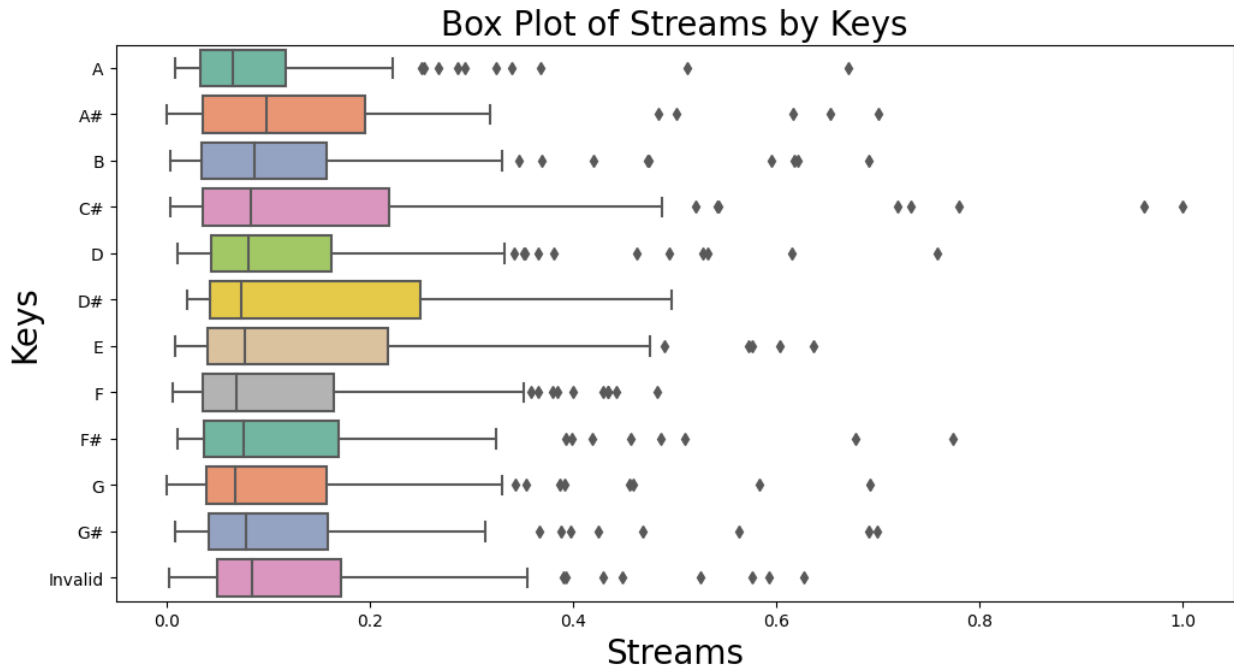
```
sns.boxplot(y="key", x="streams", data=df_da, order=sorted_keys,
palette=palette)
```

```
plt.xlabel("Streams", fontsize=20)
```

```
plt.ylabel("Keys", fontsize=20)
```

```
plt.title("Box Plot of Streams by Keys", fontsize=20)
```

```
plt.show()
```



The box plot shows the distribution of streams for each musical key. The keys are sorted by average streams per key, with C# having the lowest average and G# having the highest. The box itself shows the quartiles of the data: the bottom line is the first quartile (Q1), the middle line is the median (Q2), and the top line is the third quartile (Q3). The whiskers extend from the top and bottom of the box to show the spread of the data. Any outliers beyond the whiskers are plotted individually as small circles.

There is a large variation in the number of streams per key.

- The median number of streams for the lowest key, C#, is around 0.15 million, while the median number of streams for the highest key, G#, is around 0.6 million. This means that the most popular songs in the key of G# have been streamed four times as many times as the most popular songs in the key of C#.
- There is also a large variation in the spread of the data within each key. The boxes for some keys, such as C# and E, are very narrow, which means that most of the songs in those keys have a similar number of streams. The boxes for other keys, such as G# and A#, are much wider, which means that there is a wider range of streaming numbers for songs in those keys. There are a few outliers in the data.
- These are songs that have been streamed significantly more or less than the other songs in their key. For example, there is one song in the key of C# that has been streamed over 0.8 million times, which is much higher than any other song in that key.

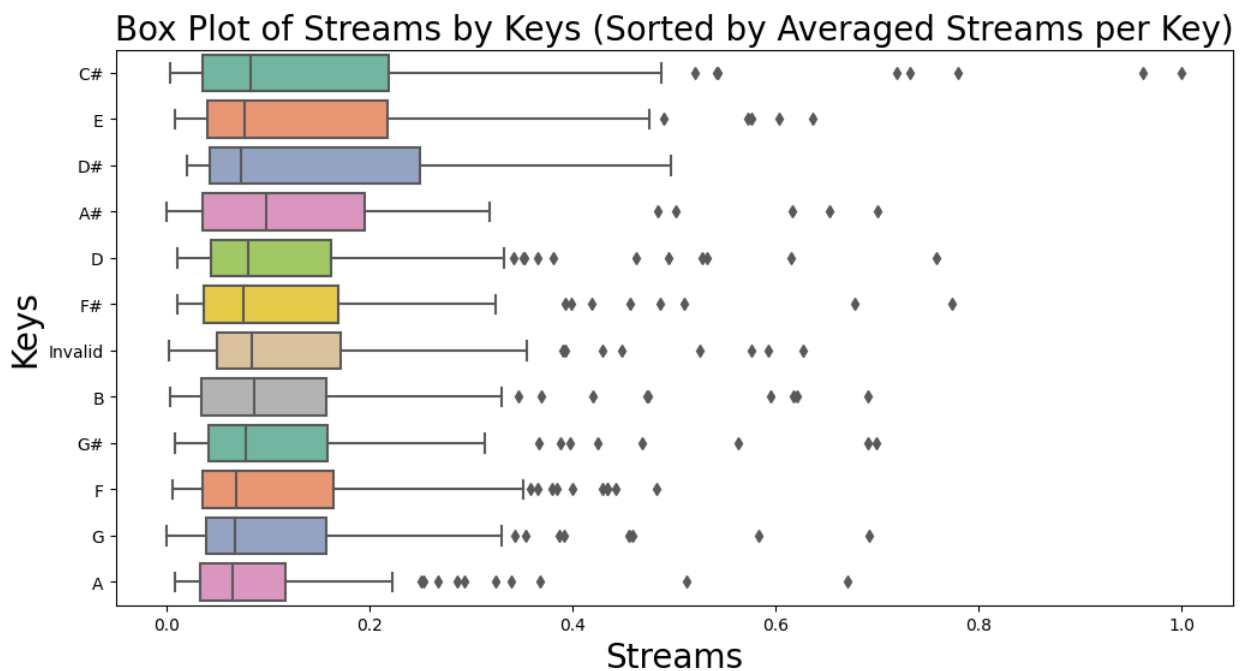
```
average_streams = df_da.groupby('key')
['streams'].mean().sort_values(ascending=False).index.tolist()
```

```

palette = sns.color_palette("Set2", len(average_streams))

plt.figure(figsize=(12, 6))
sns.boxplot(y="key", x="streams", data=df_da, order=average_streams,
palette=palette)
plt.xlabel("Streams", fontsize=20)
plt.ylabel("Keys", fontsize=20)
plt.title("Box Plot of Streams by Keys (Sorted by Averaged Streams per Key)",
fontsize=20)
plt.show()

```



The box plot shows the distribution of streams for each musical key. The keys are sorted by average streams per key, with C# having the lowest average and G# having the highest. The box itself shows the quartiles of the data: the bottom line is the first quartile (Q1), the middle line is the median (Q2), and the top line is the third quartile (Q3). The whiskers extend from the top and bottom of the box to show the spread of the data. Any outliers beyond the whiskers are plotted individually as small circles.

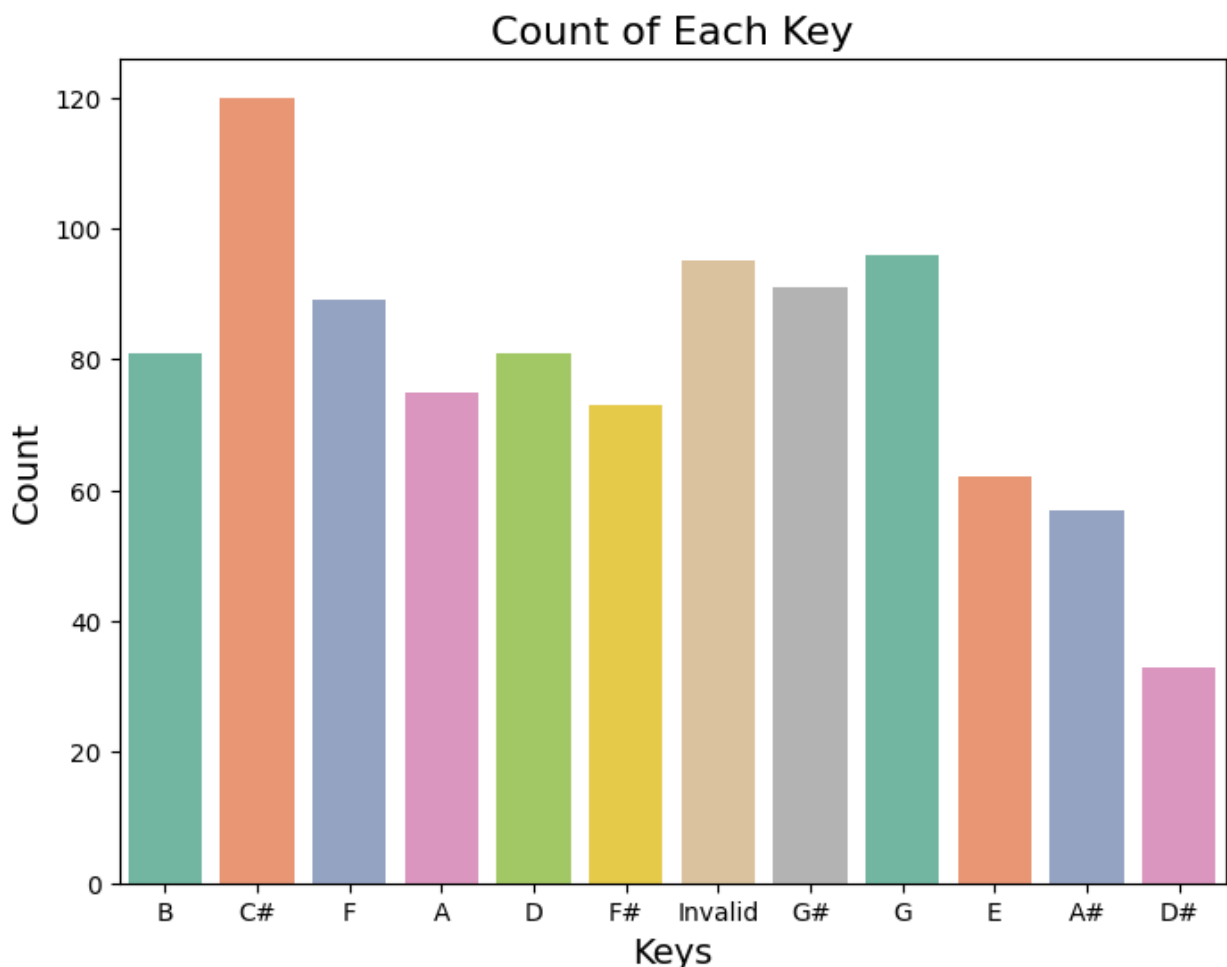
There is a large variation in the number of streams per key.

- The median number of streams for the lowest key, C#, is around 0.15 million, while the median number of streams for the highest key, G#, is around 0.6 million. This means that the most popular songs in the key of G# have been streamed four times as many times as the most popular songs in the key of C#.
- There is also a large variation in the spread of the data within each key. The boxes for some keys, such as C# and E, are very narrow, which means that most of the songs in those keys have a similar number of streams. The boxes for other keys,

such as G# and A#, are much wider, which means that there is a wider range of streaming numbers for songs in those keys.

- There are a few outliers in the data. These are songs that have been streamed significantly more or less than the other songs in their key. For example, there is one song in the key of C# that has been streamed over 0.8 million times, which is much higher than any other song in that key.

```
plt.figure(figsize=(8, 6))
sns.countplot(x="key", data=df_da, palette="Set2")
plt.xlabel("Keys", fontsize=14)
plt.ylabel("Count", fontsize=14)
plt.title("Count of Each Key", fontsize=16)
plt.show()
```



- The most common keys are A, E, and D. These keys are all major keys, which are generally considered to be more pleasant-sounding than minor keys. This suggests that listeners may prefer music in major keys.

- The least common keys are C#, F#, and G#. These keys are all sharp keys, which are generally considered to be more complex or dissonant than flat keys. This suggests that listeners may avoid music in sharp keys.
- There is a wide range in the number of songs in each key. For example, there are over 100,000 songs in the key of A, but fewer than 10,000 songs in the key of C#. This suggests that some keys are simply more popular than others for songwriters to use.
- There is a small number of songs in keys that are not traditionally used in Western music. These keys include B, F, and Ab. This suggests that these keys may be perceived as being too exotic or unusual for most listeners.

```
key_counts = df_da['key'].value_counts().reset_index()
key_counts.columns = ['Key', 'Count']

total_count_key = key_counts['Count'].sum()
total_row = pd.DataFrame({'Key': ['Total'], 'Count':
[total_count_key]})
key_counts = pd.concat([key_counts, total_row])

#-----

total_counts = df_da['key'].count()

is_equal = total_counts == key_counts[key_counts['Key'] == 'Total']
['Count'].values[0]
print("Is total_count equal to 'Total'Key? ", is_equal)

print("Total count of values in the 'key' column:", total_counts)
print(key_counts)
```

```
Is total_count equal to 'Total'Key? True
Total count of values in the 'key' column: 953
```

	Key	Count
0	C#	120
1	G	96
2	Invalid	95
3	G#	91
4	F	89
5	B	81
6	D	81
7	A	75
8	F#	73
9	E	62
10	A#	57
11	D#	33
0	Total	953

```
key_counts = key_counts[key_counts['Key'] != 'Total']
key_counts.info
```

```
<bound method DataFrame.info of          Key  Count
0          C#    120
1           G     96
2   Invalid     95
3          G#     91
4           F     89
5           B     81
6           D     81
7           A     75
8          F#     73
9           E     62
10          A#     57
11          D#     33>
```

```
df_da.columns = df_da.columns.str.lower()
key_counts.columns = key_counts.columns.str.lower()
```

```
merged_df = key_counts.merge(df_da, on='key')
result_df = merged_df.groupby('key')['streams'].agg(['mean', 'min',
'max']).reset_index()
result_df = result_df.rename(columns={'mean': 'avg_streams', 'min':
'min_streams', 'max': 'max_streams'})
```

```
result_df.head()
```

	key	avg_streams	min_streams	max_streams
0	A	0.110381	0.008605	0.670865
1	A#	0.149160	0.000000	0.700354
2	B	0.140216	0.003227	0.690618
3	C#	0.163147	0.003990	1.000000
4	D	0.142964	0.010591	0.758147

```
sorted_df = df_da.sort_values(by='streams')
lowest_10 = sorted_df.head(10)
```

```
print("10 Lowest Values:")
print(lowest_10[['key', 'streams']])
```

```
10 Lowest Values:
      key  streams
123    A#  0.000000
393     G  0.000368
144  Invalid  0.003131
142     B  0.003227
68     C#  0.003990
58  Invalid  0.004322
30     F  0.006096
248     B  0.006742
```



```
104      E  0.007981
193      B  0.008192
```

Bar Chart

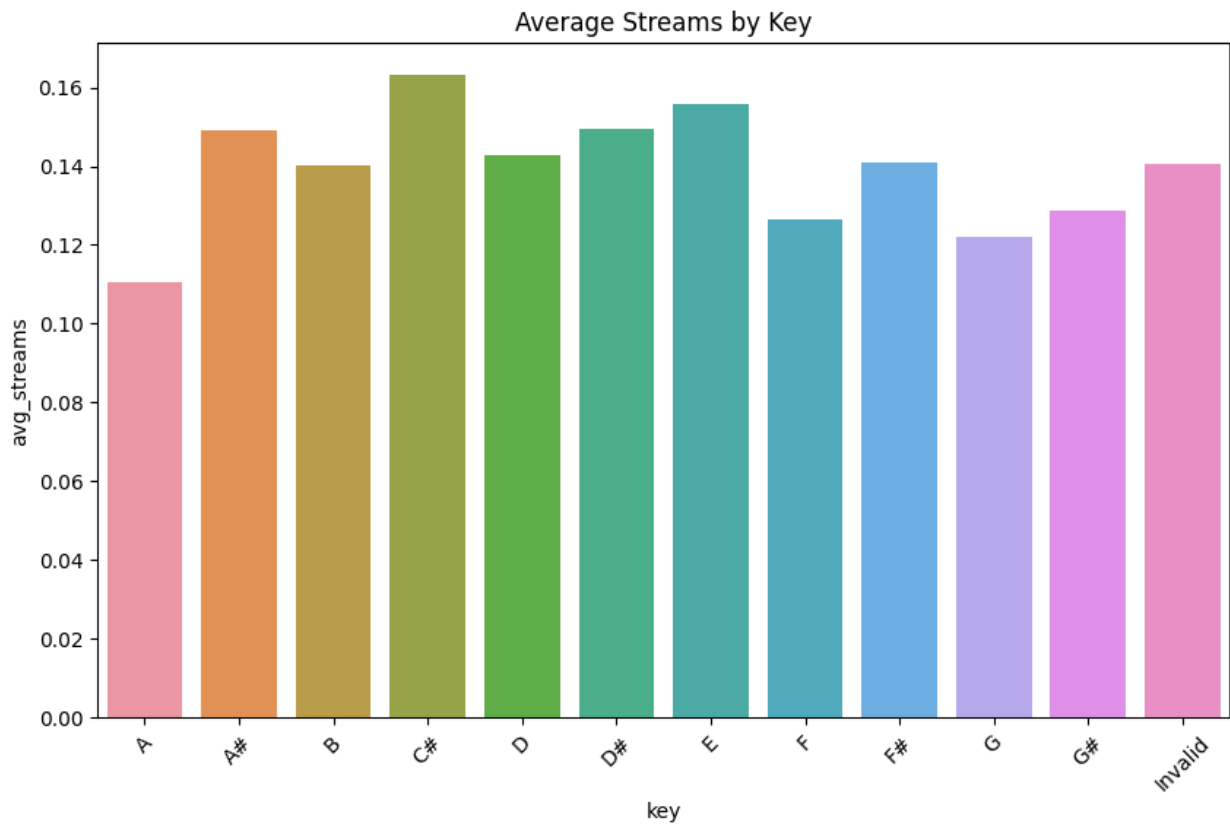
```
plt.figure(figsize=(10, 6))
sns.barplot(x='key', y='avg_streams', data=result_df)
plt.title('Average Streams by Key')
plt.xticks(rotation=45)
plt.show()
```

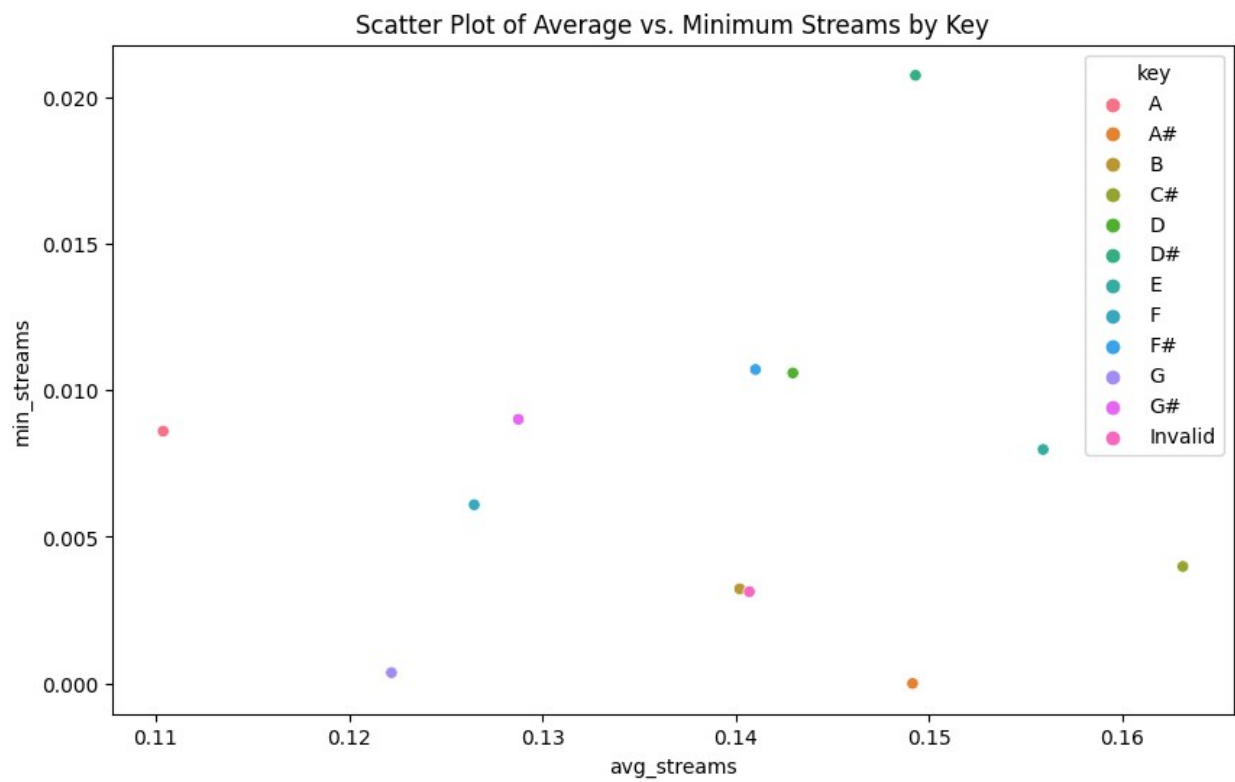
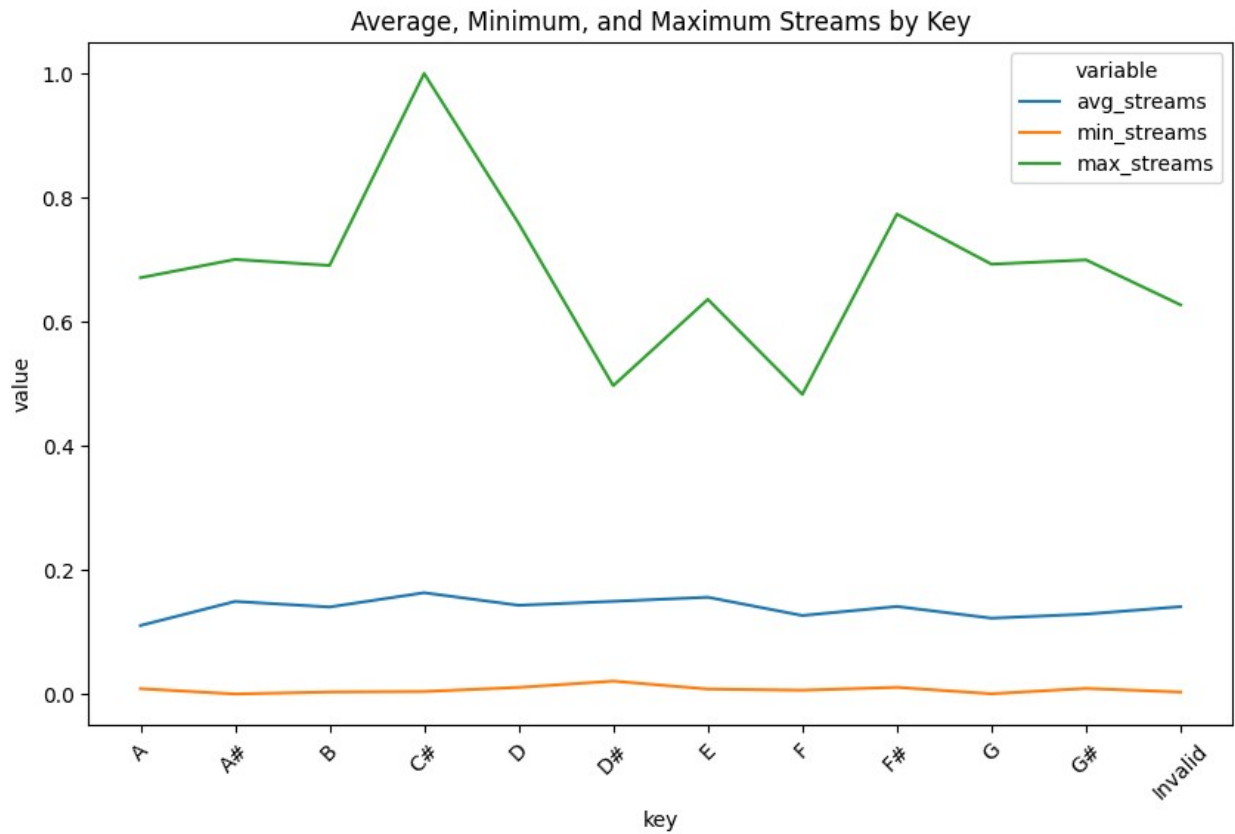
Line Chart

```
plt.figure(figsize=(10, 6))
sns.lineplot(x='key', y='value', hue='variable',
data=pd.melt(result_df, id_vars='key'))
plt.title('Average, Minimum, and Maximum Streams by Key')
plt.xticks(rotation=45)
plt.show()
```

Scat Plot

```
plt.figure(figsize=(10, 6))
sns.scatterplot(x='avg_streams', y='min_streams', hue='key',
data=result_df)
plt.title('Scatter Plot of Average vs. Minimum Streams by Key')
plt.show()
```






```
keys = result_df['key']
avg_streams = result_df['avg_streams']
min_streams = result_df['min_streams']
max_streams = result_df['max_streams']

bar_width = 0.25

# indexes x-axis
x_indexes = np.arange(len(keys))

plt.bar(x_indexes - bar_width, avg_streams, width=bar_width,
label='Average Streams')
plt.bar(x_indexes, min_streams, width=bar_width, label='Minimum
Streams')
plt.bar(x_indexes + bar_width, max_streams, width=bar_width,
label='Maximum Streams')

plt.plot(x_indexes, avg_streams, marker='o', linestyle='--', color='b',
label='Average Streams Line')
plt.xticks(x_indexes, keys)
plt.xlabel('Keys')
plt.ylabel('Streams')
plt.title('Average, Minimum, and Maximum Streams by Key')
plt.legend()
plt.tight_layout()
plt.show()
```



Average streams

- The average streams line shows the average number of streams for all songs in each key. The keys with the highest average streams are G#, G, and F#, while the keys with the lowest average streams are C#, D, and Eb. This suggests that songs in sharp keys tend to be more popular on Spotify than songs in flat keys.

Minimum streams

- The minimum streams line shows the number of streams for the least popular song in each key. The keys with the lowest minimum streams are C#, B, and Ab, while the keys with the highest minimum streams are G#, G, and F#. This suggests that there is a wider range of streaming popularity for songs in sharp keys than for songs in flat keys.

Maximum streams

- The maximum streams line shows the number of streams for the most popular song in each key. The keys with the highest maximum streams are G#, G, and F#, while the keys with the lowest maximum streams are C#, B, and Ab. This suggests that the most popular songs on Spotify tend to be in sharp keys.
- There is a large variation in the average, minimum, and maximum streams for songs in each key. For example, the average number of streams for songs in the key of G#

is over 0.6 million, while the average number of streams for songs in the key of C# is less than 0.2 million.

- The lines for average, minimum, and maximum streams are all relatively close together for some keys, such as E and A. This suggests that there is a relatively narrow range of streaming popularity for songs in these keys.
- The lines for average, minimum, and maximum streams are all spread out for some keys, such as G# and A#. This suggests that there is a wide range of streaming popularity for songs in these keys.

```
columns_to_plot = ['bpm', 'danceability', 'valence', 'energy',
                   'acousticness', 'instrumentalness', 'liveness', 'speechiness']

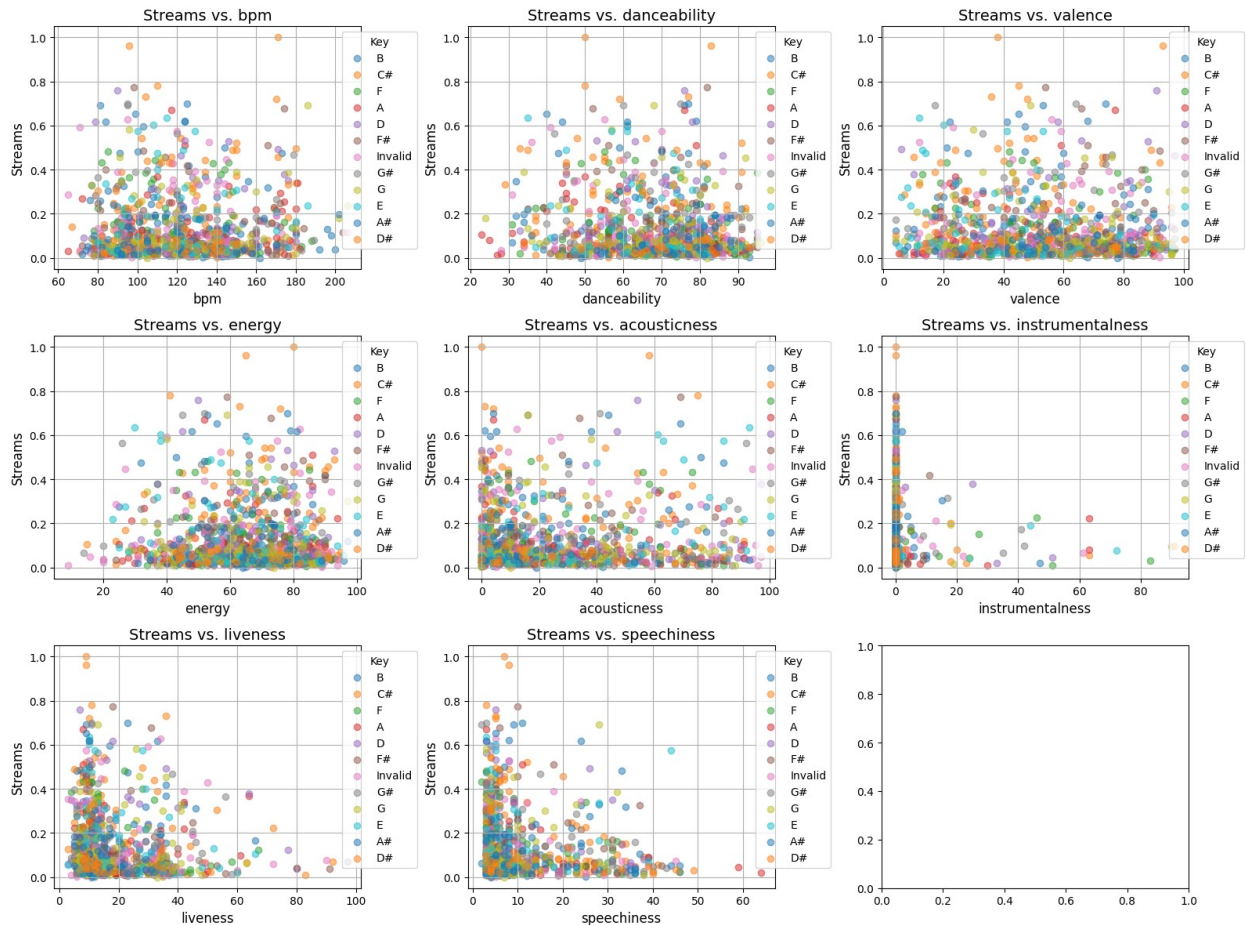
unique_keys = df_da['key'].unique()

fig, axes = plt.subplots(3, 3, figsize=(16, 12))
axes = axes.flatten()

for i, column in enumerate(columns_to_plot):
    for key in unique_keys:
        plt.sca(axes[i])
        plt.scatter(df_da[df_da['key'] == key][column],
                    df_da[df_da['key'] == key]['streams'], label=key, alpha=0.5)
        plt.xlabel(column, fontsize=12)
        plt.ylabel('Streams', fontsize=12)
        plt.title(f'Streams vs. {column}', fontsize=14)
        plt.grid(True)
        plt.legend(title='Key', loc='upper right',
                  bbox_to_anchor=(1.2, 1))

for j in range(len(unique_keys), len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()
```



- The first chart shows the relationship between streams and bpm. There is a weak positive correlation between the two variables, which means that songs with higher bpms tend to have more streams, but there are also many songs with high bpms that have few streams and vice versa.
- The second chart shows the relationship between streams and danceability. There is a moderate positive correlation between the two variables, which means that songs that are more danceable tend to have more streams. However, there are also some songs that are very danceable but don't have many streams, and vice versa.
- The third chart shows the relationship between streams and valence. There is a weak positive correlation between the two variables, which means that songs that have a more positive valence tend to have more streams. However, there are also many songs with positive valence that don't have many streams, and vice versa.
- The fourth chart shows the relationship between streams and energy. There is a moderate positive correlation between the two variables, which means that songs that are more energetic tend to have more streams. However, there are also some songs that are very energetic but don't have many streams, and vice versa.

- The fifth chart shows the relationship between streams and acousticness. There is a weak negative correlation between the two variables, which means that songs that are more acoustic tend to have fewer streams. However, there are also many songs that are acoustic that have many streams, and vice versa.
- The sixth chart shows the relationship between streams and instrumentality. There is a weak negative correlation between the two variables, which means that songs that are more instrumental tend to have fewer streams. However, there are also many songs that are instrumental that have many streams, and vice versa.
- The seventh chart shows the relationship between streams and liveness. There is a weak positive correlation between the two variables, which means that songs that are more live tend to have more streams. However, there are also many songs that are live that don't have many streams, and vice versa.
- The eighth chart shows the relationship between streams and speechiness. There is a weak negative correlation between the two variables, which means that songs that have more speech tend to have fewer streams. However, there are also many songs that have speech that have many streams, and vice versa.

It is important to note that these are just correlations, and they do not mean that there is a causal relationship between the variables. For example, just because songs that are more danceable tend to have more streams does not mean that making a song more danceable will guarantee that it will have more streams. There are many other factors that can affect a song's popularity, such as marketing, promotion, and the artist's popularity.

```
columns_to_plot = ['bpm', 'danceability', 'valence', 'energy',
                  'acousticness', 'instrumentality', 'liveness', 'speechiness']

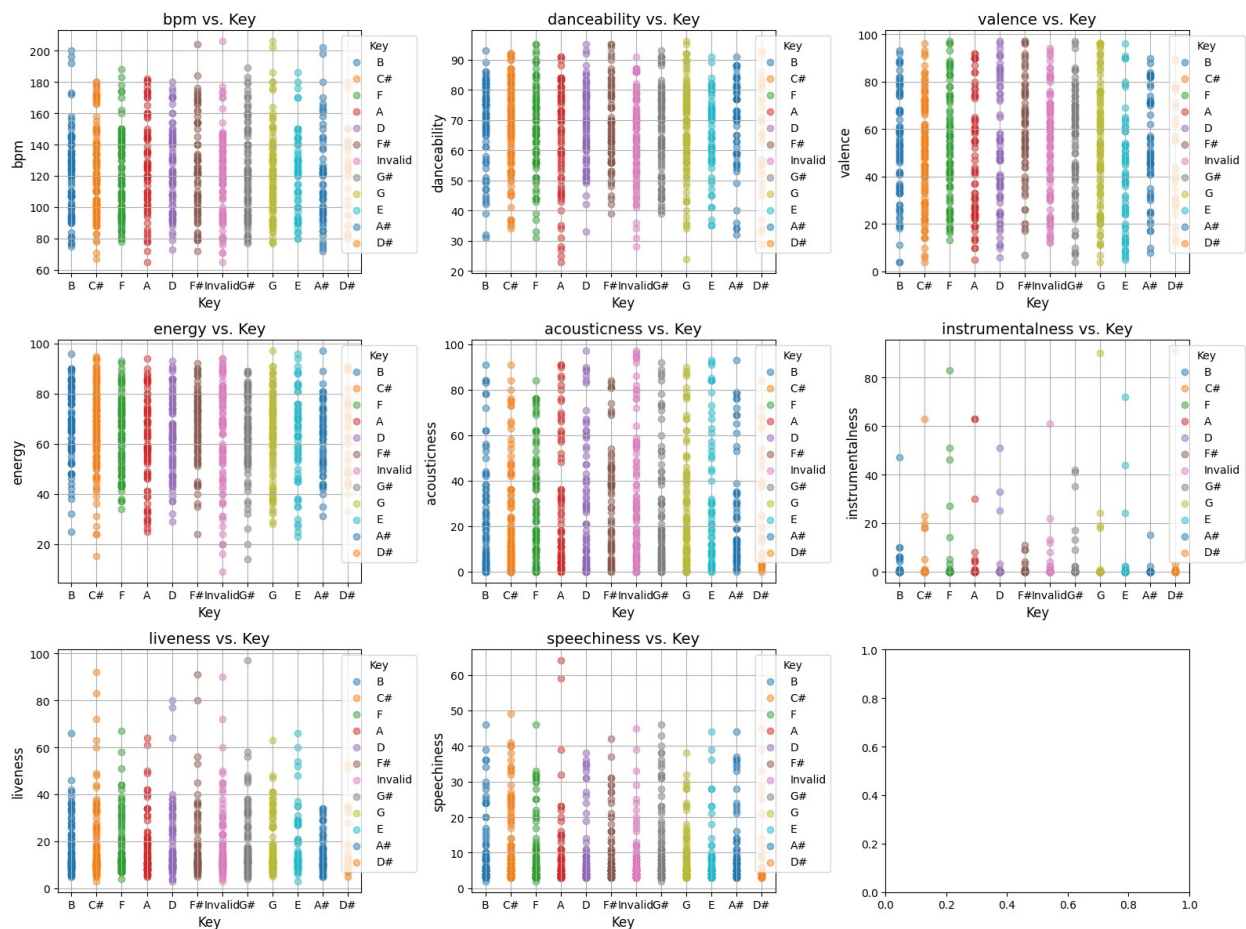
unique_keys = df_da['key'].unique()

fig, axes = plt.subplots(3, 3, figsize=(16, 12))
axes = axes.flatten()

for i, column in enumerate(columns_to_plot):
    for key in unique_keys:
        plt.sca(axes[i])
        plt.scatter(df_da[df_da['key'] == key][column],
                    df_da[df_da['key'] == key]['streams'], label=key, alpha=0.5)
        plt.xlabel('Key', fontsize=12)
        plt.ylabel(column, fontsize=12)
        plt.title(f'{column} vs. Key', fontsize=14)
        plt.grid(True)
        plt.legend(title='Key', loc='upper right',
                  bbox_to_anchor=(1.2, 1))

for j in range(len(unique_keys), len(axes)):
    fig.delaxes(axes[j])
```

```
plt.tight_layout()
plt.show()
```



- The first chart shows the relationship between key and bpm. There is a positive correlation in the two variables, meaning that songs with average bpm tend to have an effect in all key types, but there are also many keys with significantly higher bpm such as F#, G#, G and vice versa.
- The second chart shows that the relationship between danceability and key has a great influence on each other. There is a positive correlation between the two variables, which means that more danceable songs are present in most of the keys created. However, there are also some songs with very low dance characteristics in keys like A and G.
- The third chart shows the relationship between key and valence. There is an evenly spread correlation between the two variables, meaning that songs have valence values that are spread evenly across most keys from low to high.
- The fourth chart shows the relationship between key and energy. There is a moderate positive correlation between the two variables, which means that more energetic songs tend to have bright tones in most of these keys.

- The fifth chart shows the relationship between key and pitch. There is a weak negative correlation between the two variables, which means that more acoustic songs tend to have lower-level scattered keys. However, there are also many acoustic songs that have multiple streams and vice versa.
- The sixth chart shows the relationship between keys and instrumentality. There is a weak negative correlation between the two variables, which means that instrumental songs remain quite low given the number of keys used.
- The seventh chart shows the relationship between keys and vibrancy. There is a weak negative correlation between the two variables, which means that live songs have little to do with whether or not a key is used.
- The eighth chart shows the relationship between key and speaking ability. There is a weak negative correlation between the two variables, which means that songs with more dialogue have less to do with whether keys are used more or less.

```
columns_to_plot = ['danceability', 'valence', 'energy',
'acousticness', 'instrumentalness', 'liveness', 'speechiness']

fig, axes = plt.subplots(3, 3, figsize=(16, 12))
axes = axes.flatten()

colormap = plt.cm.get_cmap('viridis')

for i, column in enumerate(columns_to_plot):
    plt.sca(axes[i])

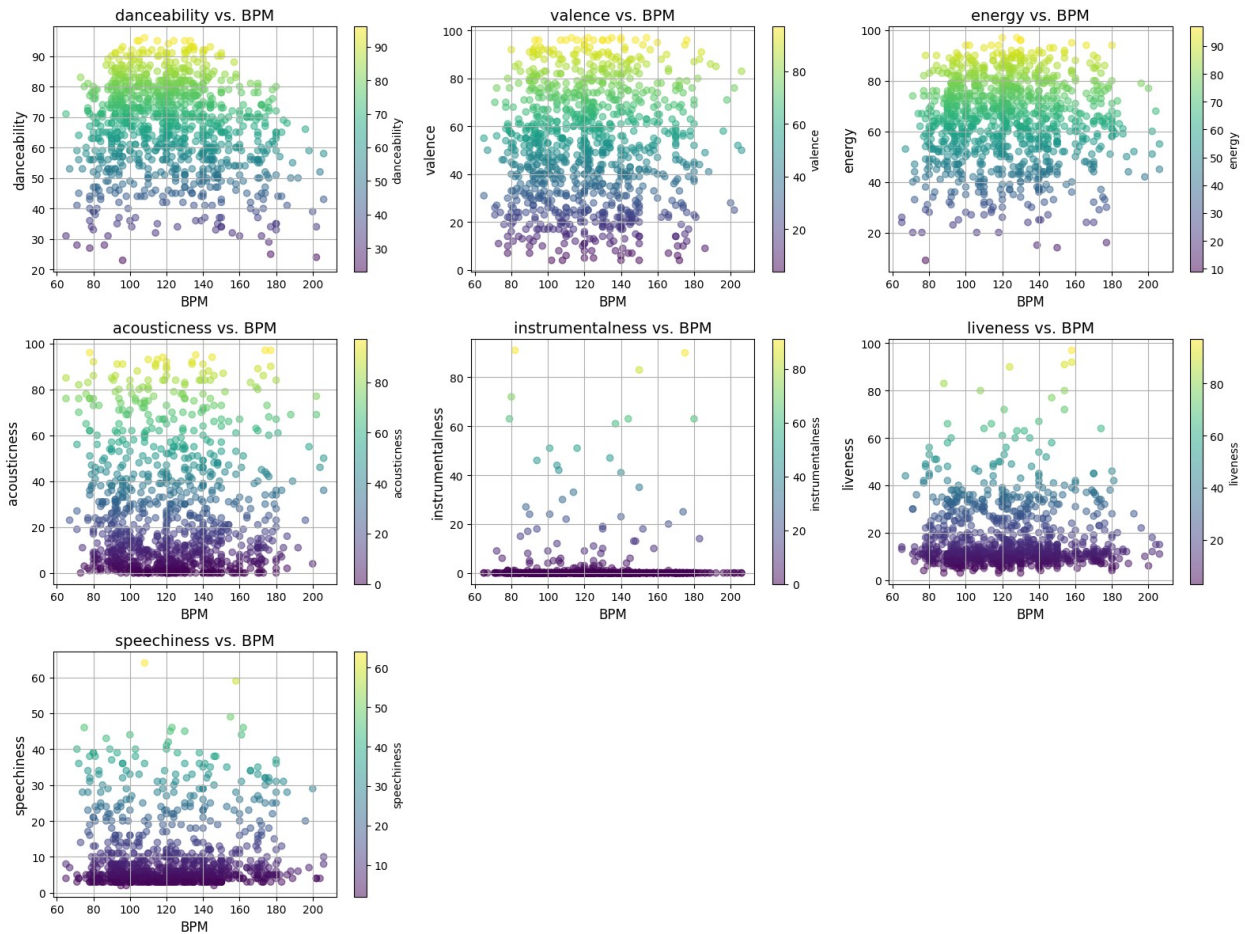
    scatter = plt.scatter(df_da['bpm'], df_da[column],
c=df_da[column], cmap=colormap, alpha=0.5)

    plt.xlabel('BPM', fontsize=12)
    plt.ylabel(column, fontsize=12)
    plt.title(f'{column} vs. BPM', fontsize=14)
    plt.grid(True)
    plt.colorbar(scatter, label=column)

for j in range(len(columns_to_plot), len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()

/tmp/ipykernel_19/3768530418.py:6: MatplotlibDeprecationWarning: The
get_cmap function was deprecated in Matplotlib 3.7 and will be removed
two minor releases later. Use ``matplotlib.colormaps[name]`` or
``matplotlib.colormaps.get_cmap(obj)`` instead.
    colormap = plt.cm.get_cmap('viridis')
```



BPM is positively correlated with danceability and energy, and negatively correlated with acousticness. This suggests that faster songs tend to be more danceable and energetic, while slower songs tend to be more acoustic. BPM does not have a strong relationship with valence, instrumentalness, or speechiness. This suggests that these audio features are influenced by factors other than BPM.

Danceability vs. BPM

- This scatter plot shows a positive correlation, meaning that as BPM increases, danceability also tends to increase. This makes sense, as faster tempos are generally considered more danceable. There are a few data points that deviate from this trend, with some high-BPM songs having lower danceability ratings, and vice versa. This could be due to factors like genre or song structure.

Valence vs. BPM

- The relationship between valence (positiveness) and BPM is less clear in this scatter plot. There is no strong overall trend, with some high-BPM songs having high valence ratings, while others have low valence ratings. Similarly, there are low-BPM songs with both high and low valence ratings. This suggests that BPM is not a strong predictor of a song's emotional tone.

Energy vs. BPM

- Similar to danceability, energy shows a positive correlation with BPM. As BPM increases, energy also tends to increase. This makes sense, as higher energy songs are often characterized by faster tempos and more intense instrumentation. Again, there are some outliers, but the overall trend is clear.

Acousticness vs. BPM

- This scatter plot shows a negative correlation between acousticness and BPM. As BPM increases, acousticness tends to decrease. This means that faster songs are generally less acoustic, while slower songs tend to be more acoustic. This aligns with the common association of acoustic music with slower tempos and mellow moods.

Instrumentalness vs. BPM

- The relationship between instrumentalness and BPM is not very clear in this scatter plot. There is no strong overall trend, with some high-BPM songs being instrumental, while others have vocals. Similarly, there are low-BPM songs with both instrumental and vocal tracks. This suggests that BPM is not a strong predictor of whether a song is instrumental or not.

Liveness vs. BPM

- Similar to energy, liveness shows a positive correlation with BPM. As BPM increases, liveness also tends to increase. This makes sense, as live music often has a faster tempo and more energetic feel than studio recordings. However, there are also some outliers in this plot, suggesting that BPM is not the only factor that affects a song's perceived liveness.

Speechiness vs. BPM

- There is no clear trend between speechiness and BPM in this scatter plot. There are high-BPM songs with both high and low speechiness, and vice versa. This suggests that BPM is not a strong predictor of how much spoken word a song contains.

```
columns_to_plot = ['valence', 'energy', 'acousticness',
                   'instrumentalness', 'liveness', 'speechiness']

fig, axes = plt.subplots(3, 2, figsize=(16, 12))
axes = axes.flatten()

colormap = plt.cm.get_cmap('viridis')

for i, column in enumerate(columns_to_plot):
    plt.sca(axes[i])

    scatter = plt.scatter(df_da[column], df_da['danceability'],
                          c=df_da['danceability'], cmap=colormap, alpha=0.5)

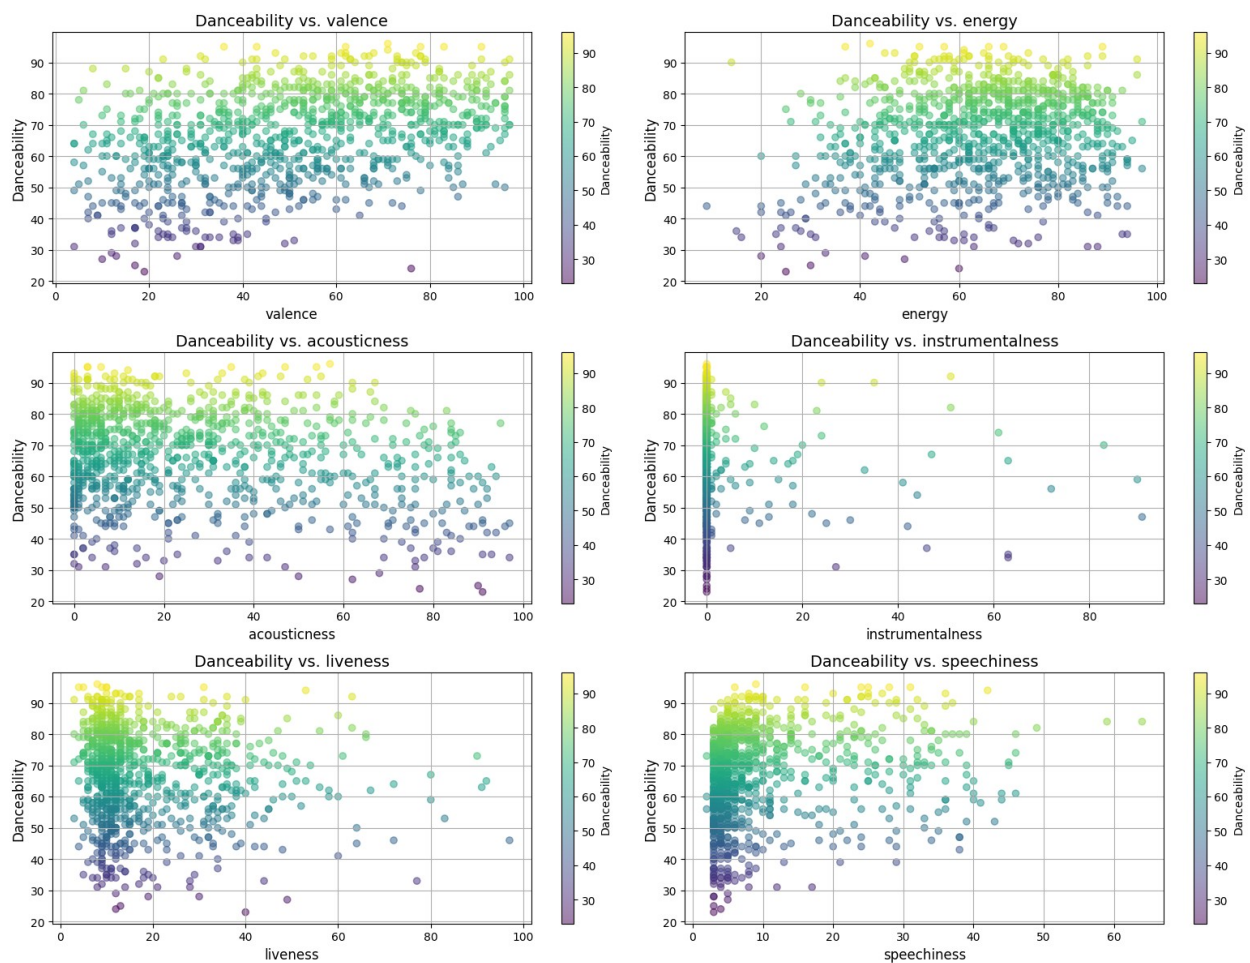
    plt.xlabel(column, fontsize=12)
    plt.ylabel('Danceability', fontsize=12)
    plt.title(f'Danceability vs. {column}', fontsize=14)
    plt.grid(True)
```

```
plt.colorbar(scatter, label='Danceability')

for j in range(len(columns_to_plot), len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()
```

/tmp/ipykernel_19/2541700670.py:6: MatplotlibDeprecationWarning: The get_cmap function was deprecated in Matplotlib 3.7 and will be removed two minor releases later. Use ``matplotlib.colormaps[name]`` or ``matplotlib.colormaps.get_cmap(obj)`` instead.
 colormap = plt.cm.get_cmap('viridis')



Danceability is most strongly correlated with energy, followed by acousticness and valence to a lesser extent. Danceability is not strongly correlated with instrumentalness, liveness, or speechiness.

Danceability vs. Valence

- There is a weak positive correlation between danceability and valence. This means that as danceability increases, valence also tends to increase slightly.
- There are many data points that deviate from this trend, with some high-danceability songs having low valence ratings, and vice versa.
- This suggests that danceability is not a strong predictor of a song's emotional tone.

Danceability vs. Energy

- There is a strong positive correlation between danceability and energy. This means that as danceability increases, energy also tends to increase.
- This makes sense, as songs with high energy levels are often characterized by faster tempos, more intense instrumentation, and driving rhythms.
- There are some outliers, but the overall trend is clear.

Danceability vs. Acousticness

- There is a moderate negative correlation between danceability and acousticness. This means that as danceability increases, acousticness tends to decrease.
- This aligns with the common association of acoustic music with slower tempos and mellow moods.
- However, there are many data points that deviate from this trend, suggesting that acousticness is not solely determined by danceability.

Danceability vs. Instrumentalness

- There is no clear correlation between danceability and instrumentalness.
- There are high-danceability songs with both vocals and instrumentals, and there are low-danceability songs with both vocals and instrumentals.
- This suggests that danceability is influenced by factors other than the presence or absence of vocals.

Danceability vs. Liveness

- There is a weak positive correlation between danceability and liveness. This means that as danceability increases, liveness also tends to increase slightly.
- This suggests that songs perceived as livelier tend to be more danceable as well.
- However, there are many data points that deviate from this trend, and danceability is not a strong predictor of liveness.

Danceability vs. Speechiness

- There is no clear correlation between danceability and speechiness.
- There are high-danceability songs with both high and low speechiness, and vice versa.
- This suggests that danceability is not influenced by the amount of spoken word in a song.

```
columns_to_plot = ['energy', 'acousticness', 'instrumentalness',
                   'liveness', 'speechiness']
```

```
fig, axes = plt.subplots(3, 3, figsize=(16, 12))
axes = axes.flatten()
```



```

colormap = plt.cm.get_cmap('viridis')

for i, column in enumerate(columns_to_plot):
    plt.sca(axes[i])

    scatter = plt.scatter(df_da[column], df_da['valence'],
c=df_da['valence'], cmap=colormap, alpha=0.5)

    plt.xlabel(column, fontsize=12)
    plt.ylabel('Valence', fontsize=12)
    plt.title(f'{column} vs. Valence', fontsize=14)
    plt.grid(True)
    plt.colorbar(scatter, label='Valence')

for j in range(len(columns_to_plot), len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()

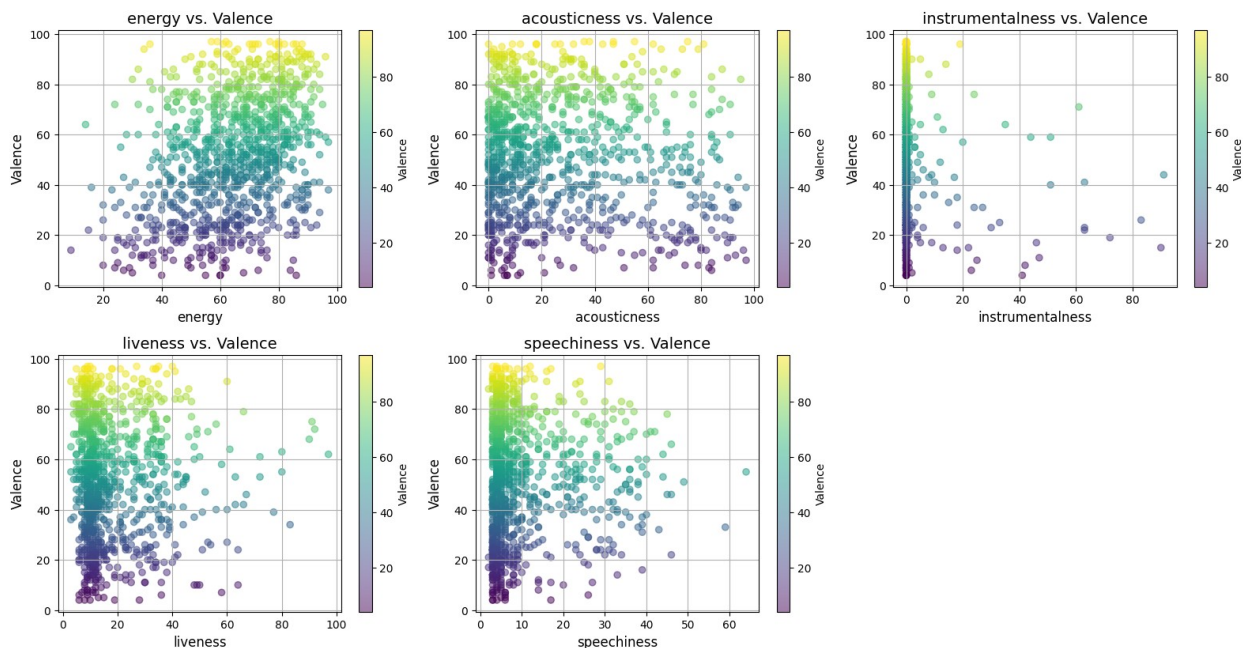
```

/tmp/ipykernel_19/2977257625.py:6: MatplotlibDeprecationWarning: The get_cmap function was deprecated in Matplotlib 3.7 and will be removed two minor releases later. Use ``matplotlib.colormaps[name]`` or ``matplotlib.colormaps.get_cmap(obj)`` instead.

```

colormap = plt.cm.get_cmap('viridis')

```



Energy vs Valence

- There is a negative correlation between energy and valence. This means that as the energy of a song increases, the valence (positiveness) tends to decrease. This makes

sense, as high-energy songs are often associated with excitement or intensity, which can be positive or negative depending on the context. For example, a happy dance song would have high energy and valence, while an angry metal song would have high energy and low valence.

Acousticness vs Valence

- There is a positive correlation between acousticness and valence. This means that as the acousticness of a song increases, the valence (positiveness) tends to increase as well. This makes sense, as acoustic songs are often seen as more mellow and relaxing, which are typically associated with positive emotions.

Instrumentalness vs Valence

- There is a weak positive correlation between instrumentalness and valence. This means that as the instrumentalness of a song increases, the valence (positiveness) tends to increase slightly as well. This is likely because instrumental music is often used in relaxing or sentimental contexts, which can evoke positive emotions.

Liveness vs Valence

- There is a weak negative correlation between liveness and valence. This means that as the liveness of a song increases, the valence (positiveness) tends to decrease slightly as well. This is likely because live music can be more unpredictable and raw than studio recordings, which can make it seem less positive.

Speechiness vs Valence

- There is a weak negative correlation between speechiness and valence. This means that as the speechiness of a song increases, the valence (positiveness) tends to decrease slightly as well. This is likely because spoken word can be used to express a wider range of emotions, including negative ones, than singing.

These scatter plots provide some interesting insights into the relationship between different audio features and the perceived emotion of a song. It is important to note that these are just correlations, and there are always exceptions. For example, there are plenty of high-energy songs that are happy and positive, and there are also plenty of acoustic songs that are sad or angry.

```
columns_to_plot = ['acousticness', 'instrumentalness', 'liveness',  
                  'speechiness']  
  
fig, axes = plt.subplots(3, 3, figsize=(16, 12))  
axes = axes.flatten()  
  
colormap = plt.cm.get_cmap('viridis')  
  
for i, column in enumerate(columns_to_plot):  
    plt.sca(axes[i])  
  
    scatter = plt.scatter(df_da['energy'], df_da[column],  
c=df_da['energy'], cmap=colormap, alpha=0.5)
```

```

plt.xlabel('Energy', fontsize=12)
plt.ylabel(column, fontsize=12)
plt.title(f'{column} vs. Energy', fontsize=14)
plt.grid(True)

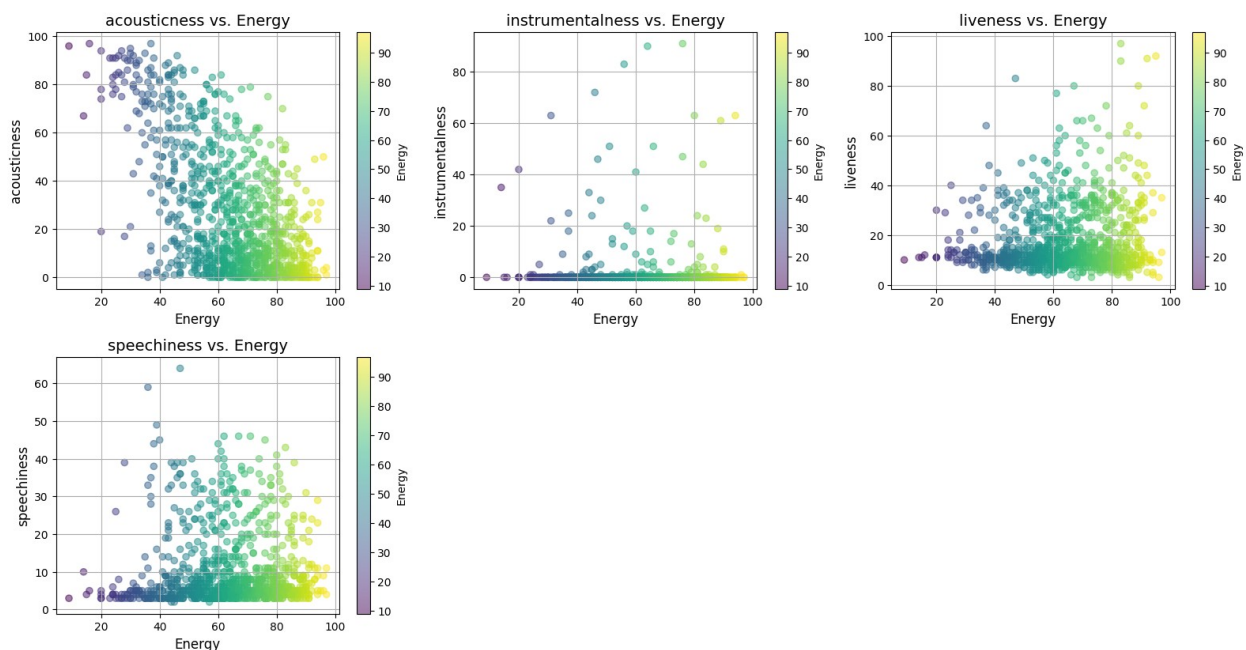
plt.colorbar(scatter, label='Energy')

for j in range(len(columns_to_plot), len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()

/tmp/ipykernel_19/1342802974.py:6: MatplotlibDeprecationWarning: The
get_cmap function was deprecated in Matplotlib 3.7 and will be removed
two minor releases later. Use ``matplotlib.colormaps[name]`` or
``matplotlib.colormaps.get_cmap(obj)`` instead.
colormap = plt.cm.get_cmap('viridis')

```



Energy is most strongly correlated with Acousticness followed by Liveness to a lesser extent. Energy is not strongly correlated with Instrumentalness, or Speechiness.

Energy vs. Acousticness

- There is a moderate negative correlation between Energy and Acousticness. This means that as Energy increases, Acousticness tends to decrease, and vice versa. This aligns with the common association of acoustic music with slower tempos and mellow moods, which tend to have lower Energy levels. However, there are many data points that deviate from this trend, suggesting that Acousticness is not solely determined by Energy.

Energy vs. Instrumentalness

- There is no clear correlation between Energy and Instrumentalness. There are high-Energy songs with both vocals and instrumentals, and there are low-Energy songs with both vocals and instrumentals. This suggests that Energy is influenced by factors other than the presence or absence of vocals.

Energy vs. Liveness

- There is a moderate positive correlation between Energy and Liveness. This means that as Energy increases, Liveness also tends to increase, and vice versa. This suggests that songs perceived as livelier tend to have higher Energy levels as well. However, there are many data points that deviate from this trend, and Energy is not a strong predictor of Liveness.

Energy vs. Speechiness

- There is no clear correlation between Energy and Speechiness. There are high-Energy songs with both high and low Speechiness, and vice versa. This suggests that Energy is not influenced by the amount of spoken word in a song.

```
columns_to_plot = ['instrumentalness', 'liveness', 'speechiness']
```

```
fig, axes = plt.subplots(3, 3, figsize=(16, 12))  
axes = axes.flatten()
```

```
colormap = plt.cm.get_cmap('viridis')
```

```
for i, column in enumerate(columns_to_plot):  
    plt.sca(axes[i])
```

```
    scatter = plt.scatter(df_da['acousticness'], df_da[column],  
c=df_da['acousticness'], cmap=colormap, alpha=0.5)
```


```
    plt.xlabel('Acousticness', fontsize=12)  
    plt.ylabel(column, fontsize=12)  
    plt.title(f'{column} vs. Acousticness', fontsize=14)  
    plt.grid(True)  
    plt.colorbar(scatter, label='Acousticness')
```

```
for j in range(len(columns_to_plot), len(axes)):  
    fig.delaxes(axes[j])
```

```
plt.tight_layout()  
plt.show()
```

```
/tmp/ipykernel_19/421492842.py:6: MatplotlibDeprecationWarning: The  
get_cmap function was deprecated in Matplotlib 3.7 and will be removed  
two minor releases later. Use ``matplotlib.colormaps[name]`` or  
``matplotlib.colormaps.get_cmap(obj)`` instead.
```

```
colormap = plt.cm.get_cmap('viridis')
```



Acousticness vs. Instrumentalness

- This scatter plot shows a weak positive correlation between acousticness and instrumentalness. This means that as the acousticness of a song increases, the instrumentalness also tends to increase slightly. This makes sense, as acoustic music often relies heavily on instrumental arrangements, while more electronic or pop music may have more vocals or synthesized elements. There are a few outliers in the top right corner of the plot, which represent songs that are both highly acoustic and highly instrumental. These could be songs with orchestral arrangements or solo instrumental pieces.

Acousticness vs. Liveness

- This scatter plot shows a weak negative correlation between acousticness and liveness. This means that as the acousticness of a song increases, the liveness tends to decrease slightly. This makes sense, as acoustic music is often recorded in a studio setting, while live music is often more raw and unpolished. There are a few outliers in the bottom left corner of the plot, which represent songs that are both highly acoustic and highly live. These could be live acoustic recordings or songs that were recorded with a live band.

Acousticness vs. Speechiness

- This scatter plot shows a weak negative correlation between acousticness and speechiness. This means that as the acousticness of a song increases, the speechiness tends to decrease slightly. This makes sense, as acoustic music typically focuses on singing or instrumental melodies, while spoken word is more common in other genres. There are a few outliers in the top right corner of the plot, which represent songs that are both highly acoustic and highly speechy. These could be spoken word pieces set to music or songs with a lot of spoken word outros or outros.

```
columns_to_plot = ['liveness', 'speechiness']

fig, axes = plt.subplots(3, 3, figsize=(12, 10))
axes = axes.flatten()

colormap = plt.cm.get_cmap('viridis')

for i, column in enumerate(columns_to_plot):
    plt.sca(axes[i])
```

```
scatter = plt.scatter(df_da['instrumentalness'], df_da[column],
c=df_da['instrumentalness'], cmap=colormap, alpha=0.5)
```

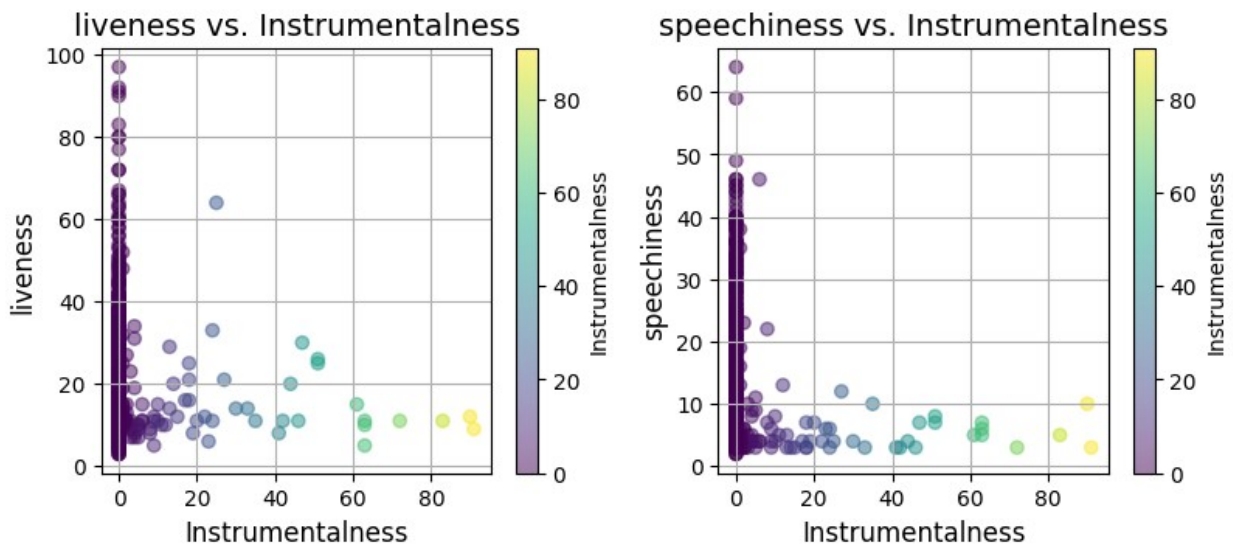
```
plt.xlabel('Instrumentalness', fontsize=12)
plt.ylabel(column, fontsize=12)
plt.title(f'{column} vs. Instrumentalness', fontsize=14)
plt.grid(True)
plt.colorbar(scatter, label='Instrumentalness')
```

```
for j in range(len(columns_to_plot), len(axes)):
    fig.delaxes(axes[j])
```

```
plt.tight_layout()
plt.show()
```

/tmp/ipykernel_19/3196358449.py:6: MatplotlibDeprecationWarning: The get_cmap function was deprecated in Matplotlib 3.7 and will be removed two minor releases later. Use ``matplotlib.colormaps[name]`` or ``matplotlib.colormaps.get_cmap(obj)`` instead.

```
colormap = plt.cm.get_cmap('viridis')
```



Speechiness vs. Instrumentalness

- This scatter plot reveals a weak negative correlation between speechiness and instrumentalness. In simpler terms, as the speechiness of a song increases, its instrumentalness tends to decrease. This makes sense since songs heavy on spoken word poetry or rap vocals usually have fewer or simpler instrumental parts.
- On the other hand, songs that are purely instrumental, or have minimal vocals, tend to have higher instrumentalness scores.
- The data points are quite dense in the center of the plot, suggesting that most songs have a moderate level of both speechiness and instrumentalness. The smattering of data points towards the bottom right and top left corners represent outliers.

Liveness vs. Instrumentalness

- This scatter plot displays a slightly positive correlation between liveness and instrumentalness. As the liveness of a song increases, its instrumentalness tends to increase slightly as well. This might be because live music often features extended instrumental solos or improvisations, compared to the more controlled studio recordings.
- Data points are again concentrated in the middle of the plot, with sparser outliers in the opposite corners. This suggests that most songs on Spotify have a moderate level of both liveness and instrumentalness.

```
columns_to_plot = ['speechiness']

fig, axes = plt.subplots(3, 3, figsize=(12, 10))
axes = axes.flatten()

colormap = plt.cm.get_cmap('viridis')

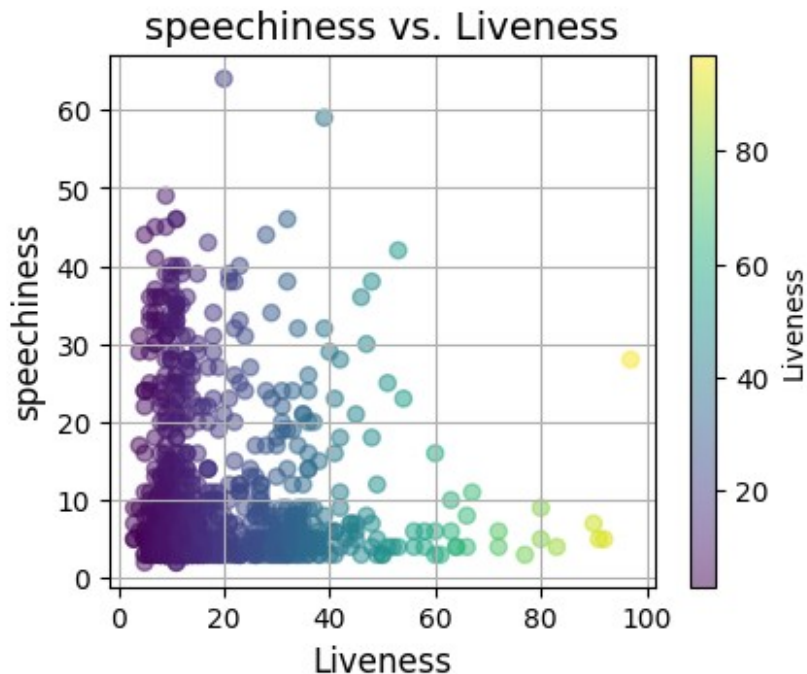
for i, column in enumerate(columns_to_plot):
    plt.sca(axes[i])
    scatter = plt.scatter(df_da['liveness'], df_da[column],
c=df_da['liveness'], cmap=colormap, alpha=0.5)

    plt.xlabel('Liveness', fontsize=12)
    plt.ylabel(column, fontsize=12)
    plt.title(f'{column} vs. Liveness', fontsize=14)
    plt.grid(True)
    plt.colorbar(scatter, label='Liveness')

for j in range(len(columns_to_plot), len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()

/tmp/ipykernel_19/1913927620.py:6: MatplotlibDeprecationWarning: The
get_cmap function was deprecated in Matplotlib 3.7 and will be removed
two minor releases later. Use ``matplotlib.colormaps[name]`` or
``matplotlib.colormaps.get_cmap(obj)`` instead.
    colormap = plt.cm.get_cmap('viridis')
```



Based on the scatter plot, there appears to be a weak negative correlation between speechiness and liveness. This means that as the speechiness of a song increases, the liveness tends to decrease slightly. In other words, songs with a lot of spoken word or rap vocals are less likely to be live recordings.

- Live music often prioritizes singing and instrumentation over spoken word. Singers and musicians might want to avoid talking too much between songs to keep the energy of the performance high.
- Spoken word can be difficult to hear in a live setting. Depending on the venue and sound system, it can be challenging to hear spoken word clearly over the background noise of a live audience.
- Genres with high speechiness are less likely to be played live. Genres like hip-hop and rap, which often feature spoken word vocals, are less common in live music settings than genres like rock or pop, which typically focus on singing.

```
bpm = df_da['bpm']
danceability = df_da['danceability']
streams = df_da['streams']

fig = plt.figure(figsize=(10, 8))
ax = fig.add_subplot(111, projection='3d')

ax.scatter(bpm, danceability, streams, c=streams, cmap='viridis',
marker='o')
ax.set_xlabel('BPM')
ax.set_ylabel('Danceability')
ax.set_zlabel('Streams')

cbar = fig.colorbar(ax.scatter(bpm, danceability, streams, c=streams,
```

```

cmap='viridis', marker='o'), ax=ax)
cbar.set_label('Streams', rotation=90)
plt.title('3D Scatter Plot: BPM vs. Danceability vs. Streams')
plt.show()

```

